

1) Problem statement.

1.This dataset comprises 244 instances that regroup a data of two regions of Algeria, namely the Bejaia region located in the northeast of Algeria and the Sidi Bel-abbes region located in the northwest of Algeria. 2.Can get insights from the dataset to know which features have contributed more in predicting Forest fire

2) Data Collection.

1.Dataset used in this particular problem statements was a dataset on Algerian Forest Fires
2.The Dataset contain data from June 2012 to September 2012. 3.Each dataset contain 122 rows and 14 columns

#Importing Pandas, Numpy, Matplotlib, Seaborn and Warnings Library.

```
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
import plotly.express as px
import warnings
```

```
warnings.filterwarnings("ignore")
```

```
%matplotlib inline
```

import dataset

```
df=pd.read_csv("C:\Algerian_forest_fires_dataset_UPDATE.csv")
```

Feature Information

Algerian Forest Fires:

1. Date : (DD/MM/YYYY) Day, month ('june' to 'september'), year (2012)
2. Temp : temperature noon (temperature max) in Celsius degrees: 22 to 42
3. RH : Relative Humidity in %: 21 to 90
4. Ws : Wind speed in km/h: 6 to 29
5. Rain: total day in mm: 0 to 16.8 FWI

Components:

1. Fine Fuel Moisture Code (FFMC) index from the FWI system: 28.6 to 92.5
2. Duff Moisture Code (DMC) index from the FWI system: 1.1 to 65.9
3. Drought Code (DC) index from the FWI system: 7 to 220.4
4. Initial Spread Index (ISI) index from the FWI system: 0 to 18.5
5. Buildup Index (BUI) index from the FWI system: 1.1 to 68
6. Fire Weather Index (FWI) Index: 0 to 31.1
7. Classes: two classes, namely Fire and not Fire

Show Top 5 Records

```
df.head()
```

	day	month	year	Temperature	RH	Ws	Rain	FFMC	DMC	DC	ISI	BUI	
FWI \	0	1	6	2012	29	57	18	0	65.7	3.4	7.6	1.3	3.4
0.5	1	2	6	2012	29	61	13	1.3	64.4	4.1	7.6	1	3.9
0.4	2	3	6	2012	26	82	22	13.1	47.1	2.5	7.1	0.3	2.7
0.1	3	4	6	2012	25	89	13	2.5	28.6	1.3	6.9	0	1.7
0	4	5	6	2012	27	77	16	0	64.8	3	14.2	1.2	3.9
0.5													

```
Classes
0 not fire
1 not fire
2 not fire
3 not fire
4 not fire
```

Shape of the dataset

```
df.shape
```

```
(247, 14)
```

Summary of the dataset

```
df.describe()
```

	day	month	year	Temperature	RH	Ws	Rain	FFMC	DMC	DC
ISI	246	245	245	245	245	245	245	245	245	245
BUI \	245	245								
count	246	245	245	245	245	245	245	245	245	245
unique	33	5	2	20	63	19	40	174	167	199

```

107 175
top      1      7 2012      35 64 14      0 88.9 7.9 8
1.1      3
freq      8     62 244      29 10 43 133      8 5 5
8      5

```

```

FWI Classes
count      245      244
unique     127      9
top        0.4  fire
freq       12     131

```

Check Datatypes in the dataset

```
df.info()
```

```

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 247 entries, 0 to 246
Data columns (total 14 columns):
#   Column                Non-Null Count  Dtype
---  -
0   day                    246 non-null   object
1   month                  245 non-null   object
2   year                   245 non-null   object
3   Temperature            245 non-null   object
4   RH                     245 non-null   object
5   Ws                     245 non-null   object
6   Rain                   245 non-null   object
7   FFMC                   245 non-null   object
8   DMC                    245 non-null   object
9   DC                     245 non-null   object
10  ISI                    245 non-null   object
11  BUI                    245 non-null   object
12  FWI                    245 non-null   object
13  Classes                244 non-null   object
dtypes: object(14)
memory usage: 27.1+ KB

```

```
#adding the region
```

```
df['Region']=[1 if i<121 else 0 for i in df.index]
```

```
df.head()
```

```

   day month  year Temperature  RH  Ws Rain  FFMC  DMC  DC  ISI  BUI
FWI \
0    1     6  2012      29  57  18     0  65.7  3.4  7.6  1.3  3.4
0.5
1    2     6  2012      29  61  13    1.3  64.4  4.1  7.6    1  3.9
0.4
2    3     6  2012      26  82  22   13.1  47.1  2.5  7.1  0.3  2.7
0.1

```

```

3      4      6      2012      25      89      13      2.5      28.6      1.3      6.9      0      1.7
0
4      5      6      2012      27      77      16      0      64.8      3      14.2      1.2      3.9
0.5

```

```

      Classes      Region
0  not fire          1
1  not fire          1
2  not fire          1
3  not fire          1
4  not fire          1

```

```

#dropping the year column
df.drop(['year'],axis=1,inplace=True)

```

```
df
```

```

      day month Temperature  RH  Ws Rain  FFMC  DMC  DC  ISI  BUI
FWI \
0      1      6          29  57  18      0  65.7  3.4  7.6  1.3  3.4
0.5
1      2      6          29  61  13      1.3  64.4  4.1  7.6  1  3.9
0.4
2      3      6          26  82  22     13.1  47.1  2.5  7.1  0.3  2.7
0.1
3      4      6          25  89  13      2.5  28.6  1.3  6.9  0  1.7
0
4      5      6          27  77  16      0  64.8  3  14.2  1.2  3.9
0.5
..      ..      ...      ...  ..  ..      ...      ...      ...      ...      ...
..
242    26      9          30  65  14      0  85.4  16  44.5  4.5  16.9
6.5
243    27      9          28  87  15      4.4  41.1  6.5  8  0.1  6.2
0
244    28      9          27  87  29      0.5  45.9  3.5  7.9  0.4  3.4
0.2
245    29      9          24  54  18      0.1  79.7  4.3  15.2  1.7  5.1
0.7
246    30      9          24  64  15      0.2  67.3  3.8  16.5  1.2  4.8
0.5

```

```

      Classes      Region
0  not fire          1
1  not fire          1
2  not fire          1
3  not fire          1
4  not fire          1
..      ..      ...      ...
242      fire          0
243  not fire          0

```

```

244    not fire          0
245    not fire          0
246    not fire          0

```

```
[247 rows x 14 columns]
```

Data cleaning

```
df.isnull().sum()
```

```

day          1
month        2
Temperature  2
RH           2
Ws           2
Rain         2
FFMC         2
DMC          2
DC           2
ISI          2
BUI          2
FWI          2
Classes      3
Region       0
dtype: int64

```

```
df[df.isnull().any(axis=1)]
```

		day	month	Temperature	RH	Ws	Rain
FFMC	\						
122		NaN	NaN	NaN	NaN	NaN	NaN
NaN							
123	Sidi-Bel Abbes Region Dataset		NaN	NaN	NaN	NaN	NaN
NaN							
168		14	7	37	37	18	0.2
88.9							

	DMC	DC	ISI	BUI	FWI	Classes	Region
122	NaN	NaN	NaN	NaN	NaN	NaN	0
123	NaN	NaN	NaN	NaN	NaN	NaN	0
168	12.9	14.6	9	12.5	10.4	fire	NaN

```

df.loc[:, 'Region']=1
df.loc[122:, 'Region']=2
df[['Region']] = df[['Region']].astype(int)

```

```
# Remove null or na values rows
```

```

df =df.dropna().reset_index(drop=True)
df.shape

```

```
(244, 14)
```

```
# Column which has string
```

```
df.iloc[[122]]
```

```
      day month Temperature  RH  Ws  Rain  FFMC  DMC  DC  ISI  BUI
FWI  \
122  day month Temperature  RH  Ws  Rain  FFMC  DMC  DC  ISI  BUI
FWI
```

```
      Classes  Region
122  Classes      2
```

```
df[df.duplicated()]
```

```
Empty DataFrame
```

```
Columns: [day, month, Temperature,  RH,  Ws, Rain , FFMC, DMC, DC,
ISI, BUI, FWI, Classes , Region]
Index: []
```

```
#remove 122th column
```

```
df = df.drop(122).reset_index(drop=True)
```

```
df.isnull().sum()
```

```
day          0
month        0
Temperature  0
  RH         0
  Ws         0
Rain         0
FFMC         0
DMC          0
DC           0
ISI          0
BUI          0
FWI          0
Classes      0
Region       0
dtype: int64
```

checking data types of each column and change it correctly

```
df.dtypes
```

```
day          object
month        object
Temperature  object
  RH         object
  Ws         object
Rain         object
FFMC         object
DMC          object
```

```

DC                object
ISI               object
BUI               object
FWI               object
Classes           object
Region            int32
dtype: object

```

```
df.columns
```

```

Index(['day', 'month', 'Temperature', ' RH', ' Ws', 'Rain ', 'FFMC',
'DMC',
      'DC', 'ISI', 'BUI', 'FWI', 'Classes ', 'Region'],
      dtype='object')

```

To fix Spaces in the column names

```

df.columns = df.columns.str.strip()
df.columns

```

```

Index(['day', 'month', 'Temperature', 'RH', 'Ws', 'Rain', 'FFMC',
'DMC', 'DC',
      'ISI', 'BUI', 'FWI', 'Classes', 'Region'],
      dtype='object')

```

To change and correct data types

```
df[['day', 'month', 'Temperature', 'RH', 'Ws']] = df[['day', 'month', 'Temperature', 'RH', 'Ws']].astype('int')
```

```

objects = [features for features in df.columns if
df[features].dtypes=='O']
for i in objects:
    if i != 'Classes':
        df[i] = df[i].astype(float)

```

```
df.dtypes
```

```

day                int32
month              int32
Temperature        int32
RH                 int32
Ws                 int32
Rain               float64
FFMC               float64
DMC                float64
DC                 float64
ISI                float64
BUI                float64
FWI                float64
Classes            object

```

Region int32

dtype: object

df.astype(str)

	day	month	Temperature	RH	Ws	Rain	FFMC	DMC	DC	ISI	BUI
FWI \											
0	1	6	29	57	18	0.0	65.7	3.4	7.6	1.3	3.4
0.5											
1	2	6	29	61	13	1.3	64.4	4.1	7.6	1.0	3.9
0.4											
2	3	6	26	82	22	13.1	47.1	2.5	7.1	0.3	2.7
0.1											
3	4	6	25	89	13	2.5	28.6	1.3	6.9	0.0	1.7
0.0											
4	5	6	27	77	16	0.0	64.8	3.0	14.2	1.2	3.9
0.5											
..
...											
238	26	9	30	65	14	0.0	85.4	16.0	44.5	4.5	16.9
6.5											
239	27	9	28	87	15	4.4	41.1	6.5	8.0	0.1	6.2
0.0											
240	28	9	27	87	29	0.5	45.9	3.5	7.9	0.4	3.4
0.2											
241	29	9	24	54	18	0.1	79.7	4.3	15.2	1.7	5.1
0.7											
242	30	9	24	64	15	0.2	67.3	3.8	16.5	1.2	4.8
0.5											

	Classes	Region
0	not fire	1
1	not fire	1
2	not fire	1
3	not fire	1
4	not fire	1
..
238	fire	2
239	not fire	2
240	not fire	2
241	not fire	2
242	not fire	2

[243 rows x 14 columns]

```
df["Classes"]=df["Classes"].str.replace("not fire","1")
```

```
df["Classes"]=df["Classes"].str.replace("fire","0")
```

```
df["Classes"]=df["Classes"].astype(int)
```

df

	day	month	Temperature	RH	Ws	Rain	FFMC	DMC	DC	ISI
BUI	FWI	\								
0	1	6	29	57	18	0.0	65.7	3.4	7.6	1.3
3.4	0.5									
1	2	6	29	61	13	1.3	64.4	4.1	7.6	1.0
3.9	0.4									
2	3	6	26	82	22	13.1	47.1	2.5	7.1	0.3
2.7	0.1									
3	4	6	25	89	13	2.5	28.6	1.3	6.9	0.0
1.7	0.0									
4	5	6	27	77	16	0.0	64.8	3.0	14.2	1.2
3.9	0.5									
..
.	...									
238	26	9	30	65	14	0.0	85.4	16.0	44.5	4.5
16.9	6.5									
239	27	9	28	87	15	4.4	41.1	6.5	8.0	0.1
6.2	0.0									
240	28	9	27	87	29	0.5	45.9	3.5	7.9	0.4
3.4	0.2									
241	29	9	24	54	18	0.1	79.7	4.3	15.2	1.7
5.1	0.7									
242	30	9	24	64	15	0.2	67.3	3.8	16.5	1.2
4.8	0.5									

	Classes	Region
0	1	1
1	1	1
2	1	1
3	1	1
4	1	1
..
238	0	2
239	1	2
240	1	2
241	1	2
242	1	2

[243 rows x 14 columns]

3.Exploring data

```
# define numerical & categorical columns
categorical_features = [i for i in df.columns if df[i].dtypes ==
'object']
numeric_features = [i for i in df.columns if df[i].dtypes != 'object']
print(categorical_features)
print(numeric_features)
```

```
[]  
['day', 'month', 'Temperature', 'RH', 'Ws', 'Rain', 'FFMC', 'DMC',  
'DC', 'ISI', 'BUI', 'FWI', 'Classes', 'Region']
```

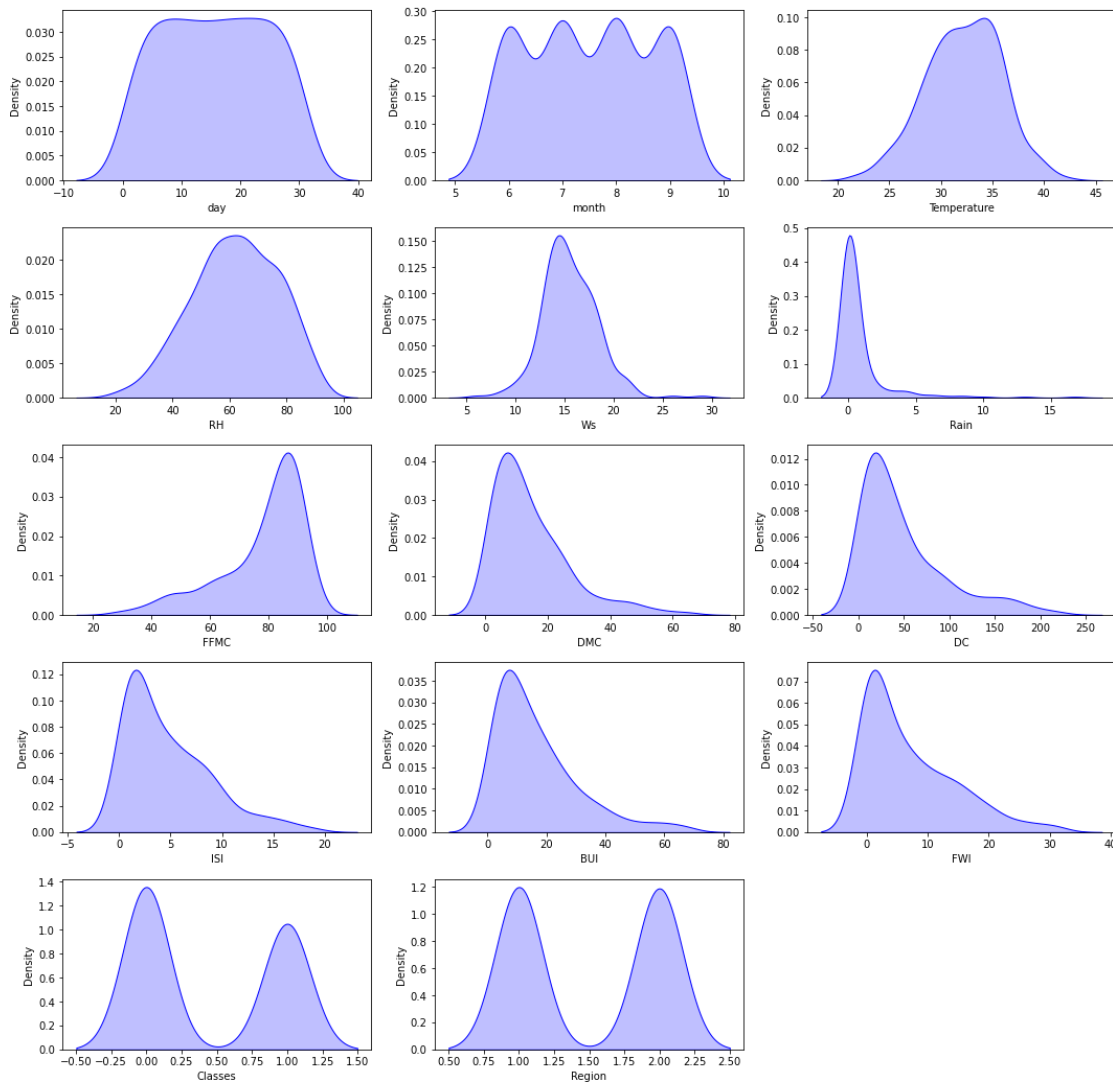
Univariate Analysis

The term univariate analysis refers to the analysis of one variable prefix “uni” means “one.” The purpose of univariate analysis is to understand the distribution of values for a single variable.

numerical columns

```
plt.figure(figsize=(15, 15))  
plt.suptitle('Univariate Analysis of Numerical Features', fontsize=20,  
fontweight='bold', alpha=0.8, y=1.)  
  
for i in range(0, len(numeric_features)):  
    plt.subplot(5, 3, i+1)  
    sns.kdeplot(x=df[numeric_features[i]], shade=True, color='b')  
    plt.xlabel(numeric_features[i])  
    plt.tight_layout()
```

Univariate Analysis of Numerical Features



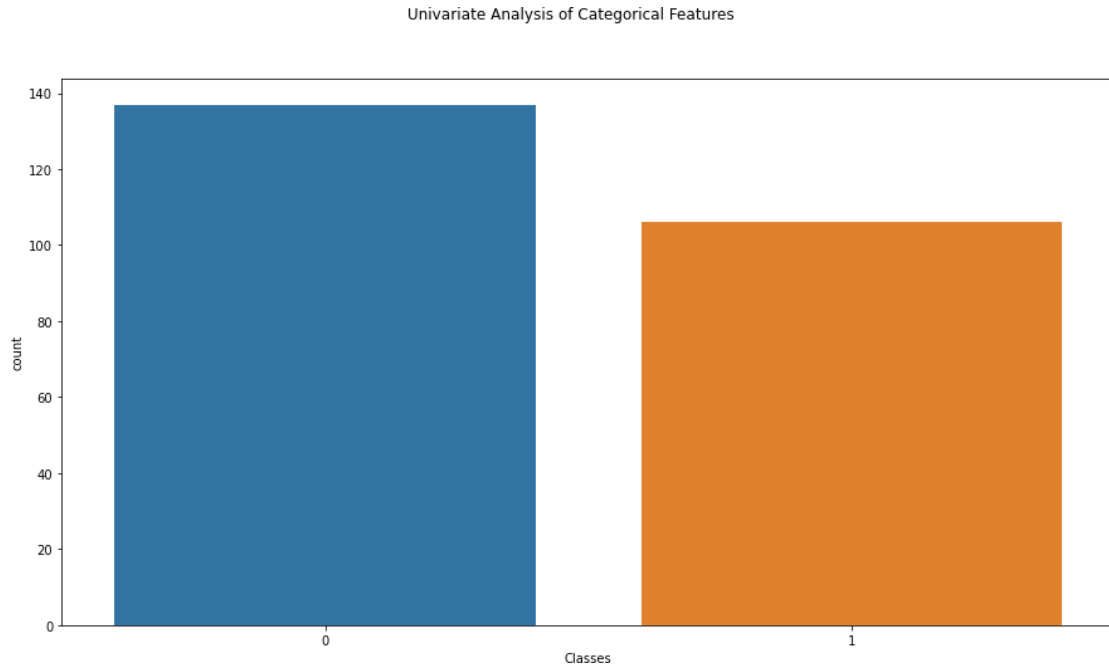
Report

- Temperature, RS, Ws are normally distributed.
- FFMC is right skewed.
- Rain, DMC, DC, ISI, BUI, FWI are left skewed or negatively skewed.

categorical analysis

```
plt.suptitle('Univariate Analysis of Categorical Features')
sns.countplot(x=df['Classes'])
```

```
<AxesSubplot:xlabel='Classes', ylabel='count'>
```




Report

from the plot we can observe that fire has occurred many times

Bivariate analysis

```
sns.relplot(x='Temperature',y='Rain',data=df,hue='Classes')
```

```
<seaborn.axisgrid.FacetGrid at 0x1754b09fa30>
```



Multivariate analysis

#Multivariate analysis is the analysis of more than two variable.

Check Multicollinearity in Numerical features

```
df.corr()
```

	day	month	Temperature	RH	Ws
Rain \					
day	1.000000	-0.000369	0.097227	-0.076034	0.047812
0.112523					
month	-0.000369	1.000000	-0.056781	-0.041252	-0.039880
0.034822					
Temperature	0.097227	-0.056781	1.000000	-0.651400	-0.284510
0.326492					
RH	-0.076034	-0.041252	-0.651400	1.000000	0.244048
0.222356					
Ws	0.047812	-0.039880	-0.284510	0.244048	1.000000
0.171506					
Rain	-0.112523	0.034822	-0.326492	0.222356	0.171506
1.000000					

FFMC	0.224956	0.017030	0.676568	-0.644873	-0.166548	-
0.543906						
DMC	0.491514	0.067943	0.485687	-0.408519	-0.000721	-
0.288773						
DC	0.527952	0.126511	0.376284	-0.226941	0.079135	-
0.298023						
ISI	0.180543	0.065608	0.603871	-0.686667	0.008532	-
0.347484						
BUI	0.517117	0.085073	0.459789	-0.353841	0.031438	-
0.299852						
FWI	0.350781	0.082639	0.566670	-0.580957	0.032368	-
0.324422						
Classes	-0.202840	-0.024004	-0.516015	0.432161	0.069964	
0.379097						
Region	0.000821	0.001857	0.269555	-0.402682	-0.181160	-
0.040013						

	FFMC	DMC	DC	ISI	BUI	
FWI \						
day	0.224956	0.491514	0.527952	0.180543	0.517117	
0.350781						
month	0.017030	0.067943	0.126511	0.065608	0.085073	
0.082639						
Temperature	0.676568	0.485687	0.376284	0.603871	0.459789	
0.566670						
RH	-0.644873	-0.408519	-0.226941	-0.686667	-0.353841	-
0.580957						
Ws	-0.166548	-0.000721	0.079135	0.008532	0.031438	
0.032368						
Rain	-0.543906	-0.288773	-0.298023	-0.347484	-0.299852	-
0.324422						
FFMC	1.000000	0.603608	0.507397	0.740007	0.592011	
0.691132						
DMC	0.603608	1.000000	0.875925	0.680454	0.982248	
0.875864						
DC	0.507397	0.875925	1.000000	0.508643	0.941988	
0.739521						
ISI	0.740007	0.680454	0.508643	1.000000	0.644093	
0.922895						
BUI	0.592011	0.982248	0.941988	0.644093	1.000000	
0.857973						
FWI	0.691132	0.875864	0.739521	0.922895	0.857973	
1.000000						
Classes	-0.769492	-0.585658	-0.511123	-0.735197	-0.586639	-
0.719216						
Region	0.222241	0.192089	-0.078734	0.263197	0.089408	
0.197102						

	Classes	Region
day	-0.202840	0.000821

```


month          -0.024004    0.001857
Temperature    -0.516015    0.269555
RH             0.432161   -0.402682
Ws             0.069964   -0.181160
Rain           0.379097   -0.040013
FFMC           -0.769492    0.222241
DMC            -0.585658    0.192089
DC             -0.511123   -0.078734
ISI            -0.735197    0.263197
BUI            -0.586639    0.089408
FWI            -0.719216    0.197102
Classes        1.000000   -0.162347
Region         -0.162347    1.000000

```

```

plt.figure(figsize = (15,10))
sns.heatmap(df.corr(), cmap="CMRmap", annot=True)
plt.show()

```



Observation


- 1.DMC-BUI,DC-BUI,FWI-BUI are very high correlated.
- 2.FFMC-BUI,ISI-BUI,DC-ISI are high correlated
- 3.Range between 0.4 to 0.6 are moderate correlated

4. correlation coefficients between 0.2 to 0.4 are less correlated
5. correlation coefficients between 0 to .1 are negatively correlated

Boxplot to find outliers in the features

```
from matplotlib import rcParams
rcParams['figure.figsize']=20,15
sns.boxplot(data=df)
```

<AxesSubplot:>




Observation

- With the help of boxplot figure we can see that some outliers are present in data


```
rcParams['figure.figsize']=15,8
sns.boxplot(x='Classes',y='Temperature',data=df)
```

<AxesSubplot:xlabel='Classes', ylabel='Temperature'>




```
from matplotlib import rcParams
rcParams['figure.figsize']=15,8
sns.boxplot(x='Classes',y='Rain',data=df)


<AxesSubplot:xlabel='Classes', ylabel='Rain'>
```



```
from matplotlib import rcParams
rcParams['figure.figsize']=15,8
p=sns.lineplot(x='Temperature',y='Rain',data=df)
```



```
rcParams['figure.figsize']=15,8
p=sns.regplot(x='Temperature',y='Rain',data=df)
```



Splitting Training and Testing data

```
#independent and dependent features
X=df.iloc[:,df.columns!='Temperature']
y=df.iloc[:,2]
```

Y

```
0    29
1    29
```

```

2      26
3      25
4      27
..
238    30
239    28
240    27
241    24
242    24

```

Name: Temperature, Length: 243, dtype: int32

X

	day	month	RH	Ws	Rain	FFMC	DMC	DC	ISI	BUI	FWI
Classes \											
0	1	6	57	18	0.0	65.7	3.4	7.6	1.3	3.4	0.5
1											
1	2	6	61	13	1.3	64.4	4.1	7.6	1.0	3.9	0.4
1											
2	3	6	82	22	13.1	47.1	2.5	7.1	0.3	2.7	0.1
1											
3	4	6	89	13	2.5	28.6	1.3	6.9	0.0	1.7	0.0
1											
4	5	6	77	16	0.0	64.8	3.0	14.2	1.2	3.9	0.5
1											
..
..											
238	26	9	65	14	0.0	85.4	16.0	44.5	4.5	16.9	6.5
0											
239	27	9	87	15	4.4	41.1	6.5	8.0	0.1	6.2	0.0
1											
240	28	9	87	29	0.5	45.9	3.5	7.9	0.4	3.4	0.2
1											
241	29	9	54	18	0.1	79.7	4.3	15.2	1.7	5.1	0.7
1											
242	30	9	64	15	0.2	67.3	3.8	16.5	1.2	4.8	0.5
1											

	Region
0	1
1	1
2	1
3	1
4	1
..	...
238	2
239	2
240	2
241	2
242	2

```
[243 rows x 13 columns]
```

```
from sklearn.model_selection import train_test_split
```

```
X_train,X_test,y_train,y_test=train_test_split(X,y,test_size=0.33,random_state=10)
```

```
X_train.head()
```

	day	month	RH	Ws	Rain	FFMC	DMC	DC	ISI	BUI	FWI
Classes \											
25	26	6	64	18	0.0	86.8	17.8	71.8	6.7	21.6	10.6
0											
121	30	9	78	14	1.4	45.0	1.9	7.5	0.2	2.4	0.1
1											
173	23	7	71	17	0.0	87.3	46.6	99.0	6.9	46.5	16.3
0											
72	12	8	51	13	0.3	81.3	15.6	75.1	2.5	20.7	4.2
1											
185	4	8	35	15	0.0	93.8	23.0	42.7	15.7	22.9	20.9
0											

	Region
25	1
121	1
173	2
72	1
185	2

```
y_train
```

25	31
121	25
173	31
72	35
185	38
	..
64	34
15	29
228	33
125	30
9	28

```
Name: Temperature, Length: 162, dtype: int32
```

```
X_train.shape
```

```
(162, 13)
```

```
y_train.shape
```

```
(162,)
```

```

# Standardization or feature scaling the dataset
from sklearn.preprocessing import StandardScaler
scaler=StandardScaler()

X_train=scaler.fit_transform(X_train)

X_test=scaler.transform(X_test) #to avoid data leakage transformer

X_train
array([[ 1.30705791, -1.39305207,  0.06835876, ...,  0.52024214,
        -0.90566259, -1.01242284],
       [ 1.77217242,  1.29354835,  0.99672801, ..., -0.93452011,
        1.10416397, -1.01242284],
       [ 0.95822202, -0.4975186 ,  0.53254338, ...,  1.30997022,
        -0.90566259,  0.9877296 ],
       ...,
       [ 0.14427163,  1.29354835, -2.45150064, ...,  1.72561657,
        -0.90566259,  0.9877296 ],
       [-1.2510719 , -1.39305207,  0.06835876, ..., -0.8098262 ,
        1.10416397,  0.9877296 ],
       [-0.55340014, -1.39305207,  1.0630401 , ..., -0.82368108,
        1.10416397, -1.01242284]])

X_test
array([[ 0.26055026, -0.4975186 ,  0.46623129, ..., -0.44959936,
        -0.90566259, -1.01242284],
       [-0.20456425,  1.29354835, -0.92632258, ...,  0.4786775 ,
        -0.90566259,  0.9877296 ],
       [ 1.77217242, -0.4975186 , -0.46213796, ...,  0.88046898,
        -0.90566259,  0.9877296 ],
       ...,
       [ 1.30705791,  0.39801488, -1.72206765, ...,  3.24965322,
        -0.90566259,  0.9877296 ],
       [-0.43712151,  1.29354835,  0.93041592, ..., -0.90681035,
        1.10416397, -1.01242284],
       [ 0.37682889,  0.39801488, -0.59476213, ...,  1.28226046,
        -0.90566259, -1.01242284]])

```

Model selection

```

from sklearn.linear_model import LinearRegression

regression = LinearRegression()

regression

LinearRegression()

regression.fit(X_train,y_train)

```

```

LinearRegression()

#print the coefficient and intercept
print(regression.coef_)

[-0.36077135 -0.2311454  -1.50045027 -0.70384333 -0.24625832
 0.84766384
 0.09164637  0.68596622  0.10380813  0.16594124 -0.38469383 -
 0.18553708
 0.22661524]

print(regression.intercept_)

32.074074074074076

#prediction for test data
reg_pred=regression.predict(X_test)

reg_pred
array([31.99493488, 33.12664016, 33.32629208, 24.7838569 ,
 29.2460991 ,
        33.66799244, 31.64931814, 34.58818421, 31.79084174,
 32.31038339,
        33.54818002, 33.34754307, 35.78255767, 32.03424221,
 34.27923643,
        33.26037242, 26.52877687, 35.94057602, 33.33186761,
 23.57110435,
        32.36221708, 32.52712236, 33.06301323, 32.81955926,
 29.903744 ,
        32.73791527, 32.89973308, 32.14375238, 31.97567765,
 34.09071758,
        34.61315534, 33.87565689, 34.55720694, 32.78051614,
 31.25884225,
        28.73433642, 32.57424997, 31.67173881, 33.05658571,
 34.00592683,
        33.83451992, 35.61866932, 34.26998499, 37.40865649,
 32.90557183,
        36.49959526, 32.27373405, 35.35128002, 30.60009684,
 31.10436839,
        32.39267139, 39.0865001 , 32.75773646, 34.64063089,
 27.20344681,
        36.86613123, 33.83117207, 33.74796958, 28.98445209,
 32.4738513 ,
        32.52599845, 31.99902312, 24.8016617 , 36.29777696,
 36.03327802,
        28.94688288, 29.45302718, 29.37313652, 35.96924316,
 28.45369633,
        29.27601583, 32.21720179, 30.29726821, 30.93046457,
 34.76079441,
        32.8057019 , 35.40714492, 31.25361667, 36.9369251 ,

```


```
26.90785041,
      33.5503991 ])
```

Assumptions of linear regression

relation between real and pred data

```
plt.scatter(y_test,reg_pred)
plt.xlabel('Test Truth Data')
plt.ylabel('Test Predict Data')

Text(0, 0.5, 'Test Predict Data')
```



#calculate residuals

```
residual=y_test-reg_pred
```

```
residual
```

```
46      -2.994935
225     -4.126640
180       2.673708
116       1.216143
124     -0.246099
```


```
...
```

```
127     -0.407145
241     -7.253617
207     -3.936925
102       3.092150
78        2.449601
```

```
Name: Temperature, Length: 81, dtype: float64
```


```
sns.displot(residual,kind='kde')
```

```
<seaborn.axisgrid.FacetGrid at 0x1754bca84c0>
```



```
plt.scatter(reg_pred, residual)
```

```
<matplotlib.collections.PathCollection at 0x1754c15d4f0>
```



Performance metrics

```
from sklearn.metrics import mean_squared_error, mean_absolute_error
print(mean_squared_error(y_test, reg_pred))
print(mean_absolute_error(y_test, reg_pred))
print(np.sqrt(mean_squared_error(y_test, reg_pred)))

7.7928584372969825
2.1532259595595913
2.7915691711467554

from sklearn.metrics import r2_score
score=r2_score(y_test, reg_pred)
print(score)

0.4370546969181548

## Adjusted R2 need to write
adjR=1-(1-score)*(len(y)-1)/(len(y)-X.shape[1]-1)
print(adjR)

0.4050971032934212
```

Ridge regression

```
from sklearn.linear_model import Ridge

ridge_reg=Ridge(alpha=1.0)

ridge_reg

Ridge()

ridge_reg.fit(X_train,y_train)

Ridge()

print(ridge_reg.coef_)

[-0.35498915 -0.22833494 -1.48029097 -0.70374866 -0.24661714
 0.85405124
 0.08174868 0.64566072 0.07203342 0.17636996 -0.314502 -
0.17870778
 0.22535392]

print(ridge_reg.intercept_)

32.074074074074076

rid_pred=ridge_reg.predict(X_test)
rid_pred

array([31.99412337, 33.12040196, 33.34904202, 24.81531826,
 29.25317586,
```

```

        33.68676155, 31.63059048, 34.61400469, 31.81080916,
32.30710445,
        33.54409223, 33.34199846, 35.75695644, 32.03039126,
34.196428 ,
        33.27080718, 26.56260115, 35.93359951, 33.31573076,
23.60085512,
        32.31428709, 32.51199227, 33.06509029, 32.82454574,
29.89856284,
        32.75584864, 32.90298721, 32.1343128 , 31.97312855,
34.06519694,
        34.59717106, 33.86491378, 34.5259007 , 32.78015204,
31.23264081,
        28.7431667 , 32.55777417, 31.67069931, 33.05264996,
33.9776409 ,
        33.85967669, 35.60302534, 34.26796346, 37.35281757,
32.90360982,
        36.53560291, 32.26187897, 35.34292679, 30.59554534,
31.10728523,
        32.39611356, 39.09427765, 32.7429986 , 34.63335351,
27.23033441,
        36.89104824, 33.84982341, 33.71470698, 29.01228005,
32.45261276,
        32.52613462, 31.98628733, 24.81143025, 36.32031857,
36.06254441,
        28.98050626, 29.42546164, 29.37614081, 35.99109657,
28.46194337,
        29.30529087, 32.22528184, 30.30676114, 30.93598706,
34.76967945,
        32.82275712, 35.37475572, 31.25349235, 36.97010006,
26.92609135,
        33.54489491])

```

```

#reation between real and predict data


```

```

plt.scatter(y_test,rid_pred)
plt.xlabel('Test Truth Data')
plt.ylabel('Test Predict Data')

Text(0, 0.5, 'Test Predict Data')

```



```
#calculate residuals
```

```
residual=y_test-rid_pred
```

```
residual
```

```
46      -2.994123
```

```
225     -4.120402
```

```
180      2.650958
```

```
116      1.184682
```

```
124     -0.253176
```

```
...
```

```
127     -0.374756
```

```
241     -7.253492
```

```
207     -3.970100
```


```
102      3.073909
```

```
78       2.455105
```

```
Name: Temperature, Length: 81, dtype: float64
```


```
sns.displot(residual,kind='kde')
```

```
<seaborn.axisgrid.FacetGrid at 0x1754cle44c0>
```



```
plt.scatter(rid_pred, residual)
```

```
<matplotlib.collections.PathCollection at 0x1754d808a60>
```



Performance metrics

```
from sklearn.metrics import mean_squared_error, mean_absolute_error
print(mean_squared_error(y_test, rid_pred))
print(mean_absolute_error(y_test, rid_pred))
print(np.sqrt(mean_squared_error(y_test, rid_pred)))

7.770134748047858
2.151505587243574
2.7874961431449297

from sklearn.metrics import r2_score
score=r2_score(y_test, rid_pred)
print(score)

0.438696224765018

## Adjusted R2 need to write
adjR=1-(1-score)*(len(y)-1)/(len(y)-X.shape[1]-1)
print(adjR)

0.4068318183106303
```

Lasso regression

```
from sklearn.linear_model import Lasso

lasso=Lasso(alpha=0.1)

lasso.fit(X_train, y_train)

Lasso(alpha=0.1)

print(lasso.coef_)

[-0.105944    -0.0832839   -1.37409593  -0.63447767  -0.1567129
 1.00625551
 0.14183148  0.32399062  0.          0.          0.          -
 0.05456468
 0.10290589]


print(lasso.intercept_)

32.074074074074076

Ls_pred=lasso.predict(X_test)

#reation between real and predict data
plt.scatter(y_test, Ls_pred)
plt.xlabel('Test Truth Data')
plt.ylabel('Test Predict Data')

Text(0, 0.5, 'Test Predict Data')
```



```
#calculate residuals
```

```
residual=y_test-Ls_pred
```

```
residual
```

```
46      -3.049143
```

```
225     -4.302261
```

```
180      2.362682
```

```
116      0.022311
```

```
124      0.240150
```

```
...
```

```
127      0.388053
```

```
241     -7.962346
```

```
207     -3.764752
```


```
102      2.464306
```

```
78       2.427680
```

```
Name: Temperature, Length: 81, dtype: float64
```


```
sns.displot(residual,kind='kde')
```

```
<seaborn.axisgrid.FacetGrid at 0x1754c193670>
```



```
plt.scatter(Ls_pred, residual)
```

```
<matplotlib.collections.PathCollection at 0x1754e194a90>
```



Performance metrics

```
from sklearn.metrics import mean_squared_error, mean_absolute_error
print(mean_squared_error(y_test, Ls_pred))
print(mean_absolute_error(y_test, Ls_pred))
print(np.sqrt(mean_squared_error(y_test, Ls_pred)))
```

```
7.615413002358264
2.1317374993041036
2.759603776334252
```

```
from sklearn.metrics import r2_score
score=r2_score(y_test, Ls_pred)
print(score)
```

```
0.4498731094372349
```

```
## Adjusted R2 need to write
adjR=1-(1-score)*(len(y)-1)/(len(y)-X.shape[1]-1)
print(adjR)
```

```
0.41864319861926136
```

Elastic net

```
from sklearn.linear_model import ElasticNet
```

```
el_reg=ElasticNet()
```

```
el_reg.fit(X_train, y_train)
```

```
ElasticNet()
```

```
print(el_reg.coef_)
```

```
[ 0.          -0.          -0.7730265  -0.26945253 -0.02396636
 0.67644044
  0.10592038  0.          0.18367796  0.03890539  0.10736141 -
 0.12263291
  0.          ]
```

```
print(el_reg.intercept_)
```

```
32.074074074074076
```

```
el_pred=el_reg.predict(X_test)
```


```
#reation between real and predict data
```

```
plt.scatter(y_test, el_pred)
```

```
plt.xlabel('Test Truth Data')
```

```
plt.ylabel('Test Predict Data')
```

```
Text(0, 0.5, 'Test Predict Data')
```





```
#calculate residuals
```

```
residual=y_test-el_pred
```


```
sns.displot(residual,kind='kde')
```

```
<seaborn.axisgrid.FacetGrid at 0x1754e21c6d0>
```



```
plt.scatter(el_pred, residual)
```

```
<matplotlib.collections.PathCollection at 0x1754e61b9a0>
```



Performance metrics

```
from sklearn.metrics import mean_squared_error, mean_absolute_error
print(mean_squared_error(y_test, el_pred))
print(mean_absolute_error(y_test, el_pred))
print(np.sqrt(mean_squared_error(y_test, el_pred)))
```

```
7.915513754032116
2.286670095063569
2.8134522839444274
```

```
from sklearn.metrics import r2_score
score=r2_score(y_test, el_pred)
print(score)
```

```
0.4281942466726337
```

```
## Adjusted R2 need to write
```

```
adjR=1-(1-score)*(len(y)-1)/(len(y)-X.shape[1]-1)
print(adjR)
```

```
0.3957336580557963
```