

Temperature Prediction for Algerian Forest Fires Dataset Data Set

Life cycle of Machine learning Project

- Understanding the Problem Statement
- Data Collection
- Exploratory data analysis
- Data Cleaning
- Data Pre-Processing
- Model Training

1) Problem statement.

- The dataset includes 244 instances that regroup a data of two regions of Algeria
- We have to predict temperature based on features

2) Data Collection.

- The Dataset is collected from uci machine learning repository (<https://archive.ics.uci.edu/ml/datasets/Algerian+Forest+Fires+Dataset++>)
- The dataset includes 244 instances that regroup a data of two regions of Algeria,namely the Bejaia region located in the northeast of Algeria and the Sidi Bel-abbes region located in the northwest of Algeria.
- 122 instances for each region.
- The dataset includes 11 attribues and 1 output attribue (class)

2.1 Import Data and Required Packages

```
In [1]: 1 import numpy as np
        2 import pandas as pd
        3 import matplotlib.pyplot as plt
        4 import seaborn as sns
```

```
In [2]: 1 df = pd.read_csv("Algerian_forest_fires_dataset_UPDATE.csv", skiprows=[0,124,125,126])
```

Attribute Information:

1. Date : (DD/MM/YYYY) Day, month ('june' to 'september'), year (2012) Weather data observations
2. Temp : temperature noon (temperature max) in Celsius degrees: 22 to 42
3. RH : Relative Humidity in %: 21 to 90
4. Ws :Wind speed in km/h: 6 to 29
5. Rain: total day in mm: 0 to 16.8 FWI Components
6. Fine Fuel Moisture Code (FFMC) index from the FWI system: 28.6 to 92.5
7. Duff Moisture Code (DMC) index from the FWI system: 1.1 to 65.9
8. Drought Code (DC) index from the FWI system: 7 to 220.4
9. Initial Spread Index (ISI) index from the FWI system: 0 to 18.5
10. Buildup Index (BUI) index from the FWI system: 1.1 to 68
11. Fire Weather Index (FWI) Index: 0 to 31.1
12. Classes: two classes, namely Fire and not Fire

Top 5 records in dataset

In [3]:

```
1 df.head()
```

Out[3]:

	day	month	year	Temperature	RH	Ws	Rain	FFMC	DMC	DC	ISI	BUI	FWI	Classes
0	1	6	2012	29	57	18	0.0	65.7	3.4	7.6	1.3	3.4	0.5	not fire
1	2	6	2012	29	61	13	1.3	64.4	4.1	7.6	1.0	3.9	0.4	not fire
2	3	6	2012	26	82	22	13.1	47.1	2.5	7.1	0.3	2.7	0.1	not fire
3	4	6	2012	25	89	13	2.5	28.6	1.3	6.9	0.0	1.7	0	not fire
4	5	6	2012	27	77	16	0.0	64.8	3.0	14.2	1.2	3.9	0.5	not fire

Last 5 records in dataset

```
In [4]: 1 df.tail()
```

Out[4]:

	day	month	year	Temperature	RH	Ws	Rain	FFMC	DMC	DC	ISI	BUI	FWI	Classes
239	26	9	2012	30	65	14	0.0	85.4	16.0	44.5	4.5	16.9	6.5	fire
240	27	9	2012	28	87	15	4.4	41.1	6.5	8	0.1	6.2	0	not fire
241	28	9	2012	27	87	29	0.5	45.9	3.5	7.9	0.4	3.4	0.2	not fire
242	29	9	2012	24	54	18	0.1	79.7	4.3	15.2	1.7	5.1	0.7	not fire
243	30	9	2012	24	64	15	0.2	67.3	3.8	16.5	1.2	4.8	0.5	not fire

Sample 5 records in dataset

```
In [5]: 1 df.sample(5)
```

Out[5]:

	day	month	year	Temperature	RH	Ws	Rain	FFMC	DMC	DC	ISI	BUI	FWI	Classes
118	27	9	2012	31	66	11	0.0	85.7	8.3	24.9	4.0	9.0	4.1	fire
14	15	6	2012	28	80	17	3.1	49.4	3.0	7.4	0.4	3.0	0.1	not fire
211	29	8	2012	35	53	17	0.5	80.2	20.7	149.2	2.7	30.6	5.9	fire
97	6	9	2012	29	74	19	0.1	75.8	3.6	32.2	2.1	5.6	0.9	not fire
21	22	6	2012	31	67	17	0.1	79.1	7.0	39.5	2.4	9.7	2.3	not fire

Datatype in dataset

```
In [6]: 1 df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 244 entries, 0 to 243
Data columns (total 14 columns):
 #   Column          Non-Null Count  Dtype
---  -
 0   day             244 non-null    int64
 1   month           244 non-null    int64
 2   year            244 non-null    int64
 3   Temperature     244 non-null    int64
 4   RH              244 non-null    int64
 5   Ws              244 non-null    int64
 6   Rain            244 non-null    float64
 7   FFMC            244 non-null    float64
 8   DMC             244 non-null    float64
 9   DC              244 non-null    object
10  ISI             244 non-null    float64
11  BUI             244 non-null    float64
12  FWI             244 non-null    object
13  Classes         243 non-null    object
dtypes: float64(5), int64(6), object(3)
memory usage: 26.8+ KB
```

Observations

1. there are missing values in Classes column
2. FWI column has data type as object
3. DC column has data type as object

3. EDA

```
In [7]: 1 df['FWI'].unique()
```

```
Out[7]: array(['0.5', '0.4', '0.1', '0', '2.5', '7.2', '7.1', '0.3', '0.9', '5.6',  
              '7.1 ', '0.2', '1.4', '2.2', '2.3', '3.8', '7.5', '8.4', '10.6',  
              '15', '13.9', '3.9', '12.9', '1.7', '4.9', '6.8', '3.2', '8',  
              '0.6', '3.4', '0.8', '3.6', '6', '10.9', '4', '8.8', '2.8', '2.1',  
              '1.3', '7.3', '15.3', '11.3', '11.9', '10.7', '15.7', '6.1', '2.6',  
              '9.9', '11.6', '12.1', '4.2', '10.2', '6.3', '14.6', '16.1',  
              '17.2', '16.8', '18.4', '20.4', '22.3', '20.9', '20.3', '13.7',  
              '13.2', '19.9', '30.2', '5.9', '7.7', '9.7', '8.3', '0.7', '4.1',  
              '1', '3.1', '1.9', '10', '16.7', '1.2', '5.3', '6.7', '9.5', '12',  
              '6.4', '5.2', '3', '9.6', '4.7', 'fire  ', '14.1', '9.1', '13',  
              '17.3', '30', '25.4', '16.3', '9', '14.5', '13.5', '19.5', '12.6',  
              '12.7', '21.6', '18.8', '10.5', '5.5', '14.8', '24', '26.3',  
              '12.2', '18.1', '24.5', '26.9', '31.1', '30.3', '26.1', '16',  
              '19.4', '2.7', '3.7', '10.3', '5.7', '9.8', '19.3', '17.5', '15.4',  
              '15.2', '6.5'], dtype=object)
```

```
In [8]: 1 # handling FWI row having alphabetical values
```

```
In [9]: 1 df[df['FWI'] == 'fire  ']
```

```
Out[9]:
```

	day	month	year	Temperature	RH	Ws	Rain	FFMC	DMC	DC	ISI	BUI	FWI	Classes
165	14	7	2012	37	37	18	0.2	88.9	12.9	14.6 9	12.5	10.4	fire	NaN

```
In [10]: 1 df.iloc[165]
```

```
Out[10]: day                14  
month                7  
year                2012  
Temperature                37  
RH                37  
Ws                18  
Rain                0.2  
FFMC                88.9  
DMC                12.9  
DC                14.6 9  
ISI                12.5  
BUI                10.4  
FWI                fire  
Classes                NaN  
Name: 165, dtype: object
```

```
In [11]: 1 # dropping row having alphabetical values
        2 df = df.drop(df.index[165])
```

```
In [12]: 1 df
```

Out[12]:

	day	month	year	Temperature	RH	Ws	Rain	FFMC	DMC	DC	ISI	BUI	FWI	Classes
0	1	6	2012	29	57	18	0.0	65.7	3.4	7.6	1.3	3.4	0.5	not fire
1	2	6	2012	29	61	13	1.3	64.4	4.1	7.6	1.0	3.9	0.4	not fire
2	3	6	2012	26	82	22	13.1	47.1	2.5	7.1	0.3	2.7	0.1	not fire
3	4	6	2012	25	89	13	2.5	28.6	1.3	6.9	0.0	1.7	0	not fire
4	5	6	2012	27	77	16	0.0	64.8	3.0	14.2	1.2	3.9	0.5	not fire
...
239	26	9	2012	30	65	14	0.0	85.4	16.0	44.5	4.5	16.9	6.5	fire
240	27	9	2012	28	87	15	4.4	41.1	6.5	8	0.1	6.2	0	not fire
241	28	9	2012	27	87	29	0.5	45.9	3.5	7.9	0.4	3.4	0.2	not fire
242	29	9	2012	24	54	18	0.1	79.7	4.3	15.2	1.7	5.1	0.7	not fire
243	30	9	2012	24	64	15	0.2	67.3	3.8	16.5	1.2	4.8	0.5	not fire

243 rows × 14 columns

```
In [13]: 1 df['FWI'].str.isnumeric().sum()
```

Out[13]: 28

```
In [14]: 1 # coverting datatype to float
        2 df['FWI'] = df['FWI'].astype('float')
```

In [15]: 1 df.info()

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 243 entries, 0 to 243
Data columns (total 14 columns):
#   Column          Non-Null Count  Dtype
---  -
0   day              243 non-null   int64
1   month            243 non-null   int64
2   year             243 non-null   int64
3   Temperature      243 non-null   int64
4   RH               243 non-null   int64
5   Ws               243 non-null   int64
6   Rain             243 non-null   float64
7   FFMC             243 non-null   float64
8   DMC              243 non-null   float64
9   DC              243 non-null   object
10  ISI              243 non-null   float64
11  BUI              243 non-null   float64
12  FWI              243 non-null   float64
13  Classes          243 non-null   object
dtypes: float64(6), int64(6), object(2)
memory usage: 28.5+ KB
```

```
In [16]: 1 df['DC'].unique()
```

```
Out[16]: array(['7.6', '7.1', '6.9', '14.2', '22.2', '30.5', '38.3', '38.8',  
              '46.3', '54.3', '61.4', '17', '7.8', '7.4', '8', '16', '27.1',  
              '31.6', '39.5', '47.7', '55.8', '63.8', '71.8', '80.3', '88.5',  
              '84.4', '92.8', '8.6', '8.3', '9.2', '18.5', '27.9', '37', '40.4',  
              '49.8', '9.3', '18.7', '27.7', '37.2', '22.9', '25.5', '34.1',  
              '43.1', '52.8', '62.1', '71.5', '79.9', '71.3', '79.7', '88.7',  
              '98.6', '108.5', '117.8', '127', '136', '145.7', '10.2', '10',  
              '19.8', '29.7', '39.1', '48.6', '47', '57', '67', '77', '75.1',  
              '85.1', '94.7', '92.5', '90.4', '100.7', '110.9', '120.9', '130.6',  
              '141.1', '151.3', '161.5', '171.3', '181.3', '190.6', '200.2',  
              '210.4', '220.4', '180.4', '8.7', '7.5', '7', '15.7', '24', '32.2',  
              '30.1', '8.4', '8.9', '16.6', '7.3', '24.3', '33.1', '41.3',  
              '49.3', '57.9', '41.4', '30.4', '15.2', '7.7', '16.3', '24.9',  
              '8.8', '8.2', '15.4', '17.6', '26.3', '28.9', '14.7', '22.5',  
              '37.8', '18.4', '25.6', '34.5', '43.3', '52.4', '36.7', '8.5',  
              '17.8', '27.3', '36.8', '46.4', '45.1', '35.4', '9.7', '9.9',  
              '9.5', '19.4', '10.4', '24.1', '42.3', '51.6', '61.1', '71',  
              '80.6', '90.1', '99', '56.6', '15.9', '19.7', '28.3', '37.6',  
              '47.2', '57.1', '67.2', '10.5', '21.4', '32.1', '42.7', '52.5',  
              '9.1', '9.8', '20.2', '30.9', '41.5', '55.5', '54.2', '65.1',  
              '76.4', '86.8', '96.8', '107', '117.1', '127.5', '137.7', '147.7',  
              '157.5', '167.2', '177.3', '166', '149.2', '159.1', '168.2',  
              '26.6', '17.7', '26.1', '25.2', '33.4', '50.2', '59.2', '63.3',  
              '77.8', '86', '88', '97.3', '106.3', '115.6', '28.1', '36.1',  
              '44.5', '7.9', '16.5'], dtype=object)
```

```
In [17]: 1 df['DC'].str.isnumeric().sum()
```

```
Out[17]: 27
```



```
In [18]: 1 df['DC'].value_counts()
```

```
Out[18]: 8      5
        7.6    4
        7.8    4
        8.4    4
        7.5    4
        ..
       92.5    1
       90.4    1
      100.7    1
      110.9    1
       16.5    1
Name: DC, Length: 197, dtype: int64
```

```
In [19]: 1 df['DC'].str.isalpha().sum()
```

```
Out[19]: 0
```

```
In [20]: 1 # coverting datatype to float
        2 df['DC'] = df['DC'].astype('float')
```

```
In [21]: 1 df.info()

<class 'pandas.core.frame.DataFrame'>
Int64Index: 243 entries, 0 to 243
Data columns (total 14 columns):
 #   Column          Non-Null Count  Dtype
---  -
 0   day             243 non-null   int64
 1   month           243 non-null   int64
 2   year            243 non-null   int64
 3   Temperature     243 non-null   int64
 4   RH              243 non-null   int64
 5   Ws              243 non-null   int64
 6   Rain            243 non-null   float64
 7   FFMFC           243 non-null   float64
 8   DMC             243 non-null   float64
 9   DC              243 non-null   float64
10  ISI             243 non-null   float64
11  BUI             243 non-null   float64
12  FWI             243 non-null   float64
13  Classes         243 non-null   object
dtypes: float64(7), int64(6), object(1)
memory usage: 28.5+ KB
```

```
In [22]: 1 # handling missing values
        2 df.isnull().sum()
```

```
Out[22]: day             0
month           0
year            0
Temperature     0
RH              0
Ws              0
Rain            0
FFMC            0
DMC             0
DC              0
ISI             0
BUI             0
FWI             0
Classes         0
dtype: int64
```

In [23]:

1df.describe()

Out[23]:

	day	month	year	Temperature	RH	Ws	Rain	FFMC	DMC	DC	ISI	BUI
count	243.000000	243.000000	243.0	243.000000	243.000000	243.000000	243.000000	243.000000	243.000000	243.000000	243.000000	243.000000
mean	15.761317	7.502058	2012.0	32.152263	62.041152	15.493827	0.762963	77.842387	14.680658	49.430864	4.742387	16.690535
std	8.842552	1.114793	0.0	3.628039	14.828160	2.811385	2.003207	14.349641	12.393040	47.665606	4.154234	14.228421
min	1.000000	6.000000	2012.0	22.000000	21.000000	6.000000	0.000000	28.600000	0.700000	6.900000	0.000000	1.100000
25%	8.000000	7.000000	2012.0	30.000000	52.500000	14.000000	0.000000	71.850000	5.800000	12.350000	1.400000	6.000000
50%	16.000000	8.000000	2012.0	32.000000	63.000000	15.000000	0.000000	83.300000	11.300000	33.100000	3.500000	12.400000
75%	23.000000	8.000000	2012.0	35.000000	73.500000	17.000000	0.500000	88.300000	20.800000	69.100000	7.250000	22.650000
max	31.000000	9.000000	2012.0	42.000000	90.000000	29.000000	16.800000	96.000000	65.900000	220.400000	19.000000	68.000000

Observations

1. year column has 0 standard deviation
2. Rain, DC column has very high max value, looks like presence of outliers

Check for duplicates

In [24]:

1df[df.duplicated()]

Out[24]:

day	month	year	Temperature	RH	Ws	Rain	FFMC	DMC	DC	ISI	BUI	FWI	Classes
-----	-------	------	-------------	----	----	------	------	-----	----	-----	-----	-----	---------

```
In [25]: 1 df.corr()
```

Out[25]:

	day	month	year	Temperature	RH	Ws	Rain	FFMC	DMC	DC	ISI	BUI	FWI
day	1.000000	-0.000369	NaN	0.097227	-0.076034	0.047812	-0.112523	0.224956	0.491514	0.527952	0.180543	0.517117	0.350781
month	-0.000369	1.000000	NaN	-0.056781	-0.041252	-0.039880	0.034822	0.017030	0.067943	0.126511	0.065608	0.085073	0.082639
year	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN
Temperature	0.097227	-0.056781	NaN	1.000000	-0.651400	-0.284510	-0.326492	0.676568	0.485687	0.376284	0.603871	0.459789	0.566670
RH	-0.076034	-0.041252	NaN	-0.651400	1.000000	0.244048	0.222356	-0.644873	-0.408519	-0.226941	-0.686667	-0.353841	-0.580957
Ws	0.047812	-0.039880	NaN	-0.284510	0.244048	1.000000	0.171506	-0.166548	-0.000721	0.079135	0.008532	0.031438	0.032368
Rain	-0.112523	0.034822	NaN	-0.326492	0.222356	0.171506	1.000000	-0.543906	-0.288773	-0.298023	-0.347484	-0.299852	-0.324422
FFMC	0.224956	0.017030	NaN	0.676568	-0.644873	-0.166548	-0.543906	1.000000	0.603608	0.507397	0.740007	0.592011	0.691132
DMC	0.491514	0.067943	NaN	0.485687	-0.408519	-0.000721	-0.288773	0.603608	1.000000	0.875925	0.680454	0.982248	0.875864
DC	0.527952	0.126511	NaN	0.376284	-0.226941	0.079135	-0.298023	0.507397	0.875925	1.000000	0.508643	0.941988	0.739521
ISI	0.180543	0.065608	NaN	0.603871	-0.686667	0.008532	-0.347484	0.740007	0.680454	0.508643	1.000000	0.644093	0.922895
BUI	0.517117	0.085073	NaN	0.459789	-0.353841	0.031438	-0.299852	0.592011	0.982248	0.941988	0.644093	1.000000	0.857973
FWI	0.350781	0.082639	NaN	0.566670	-0.580957	0.032368	-0.324422	0.691132	0.875864	0.739521	0.922895	0.857973	1.000000

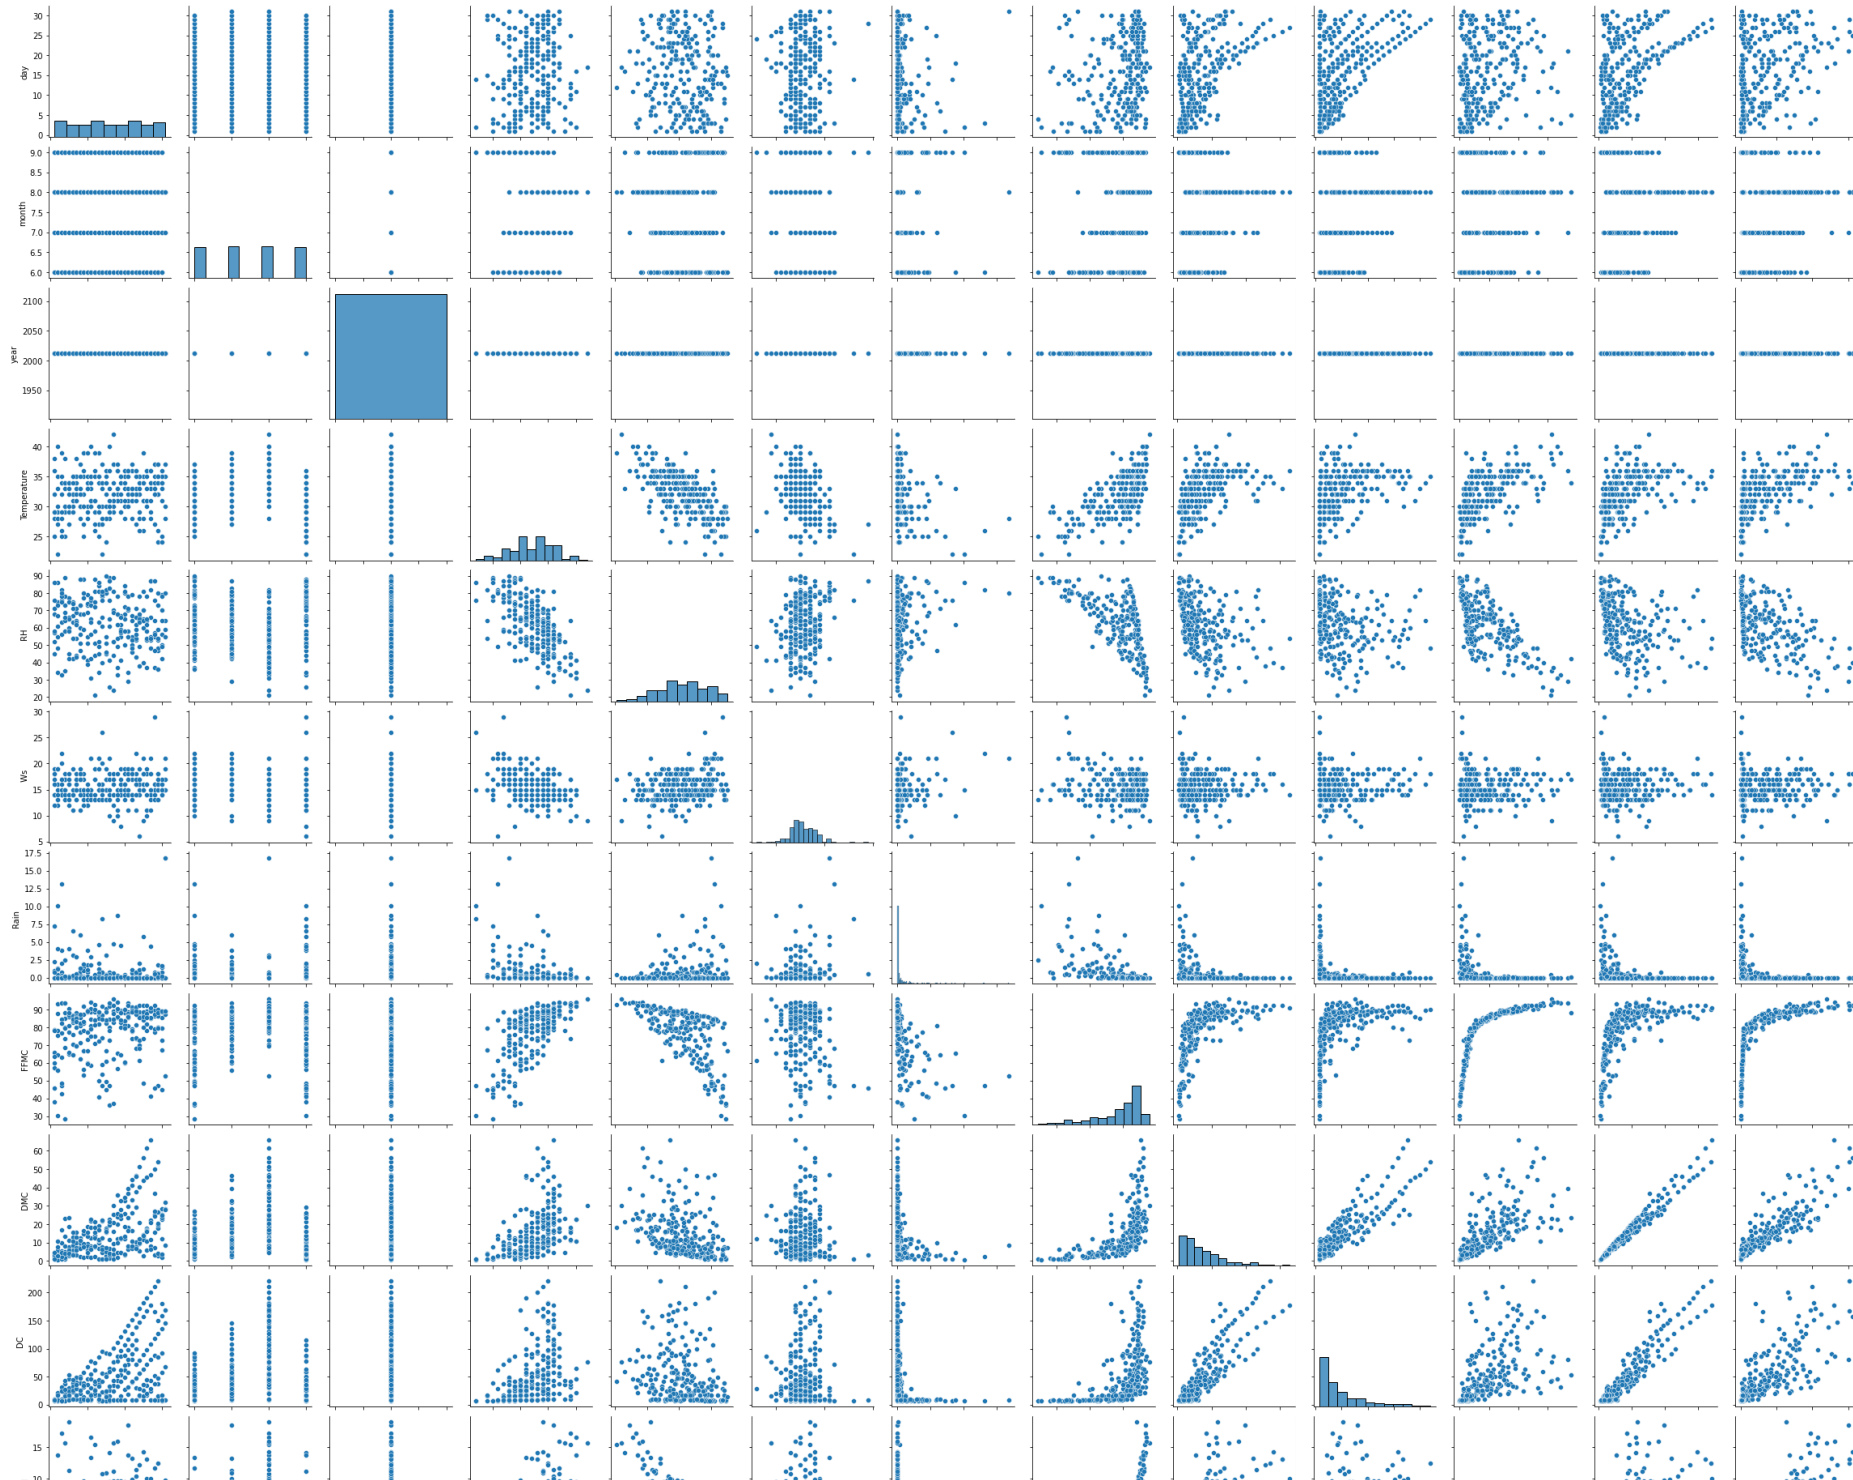


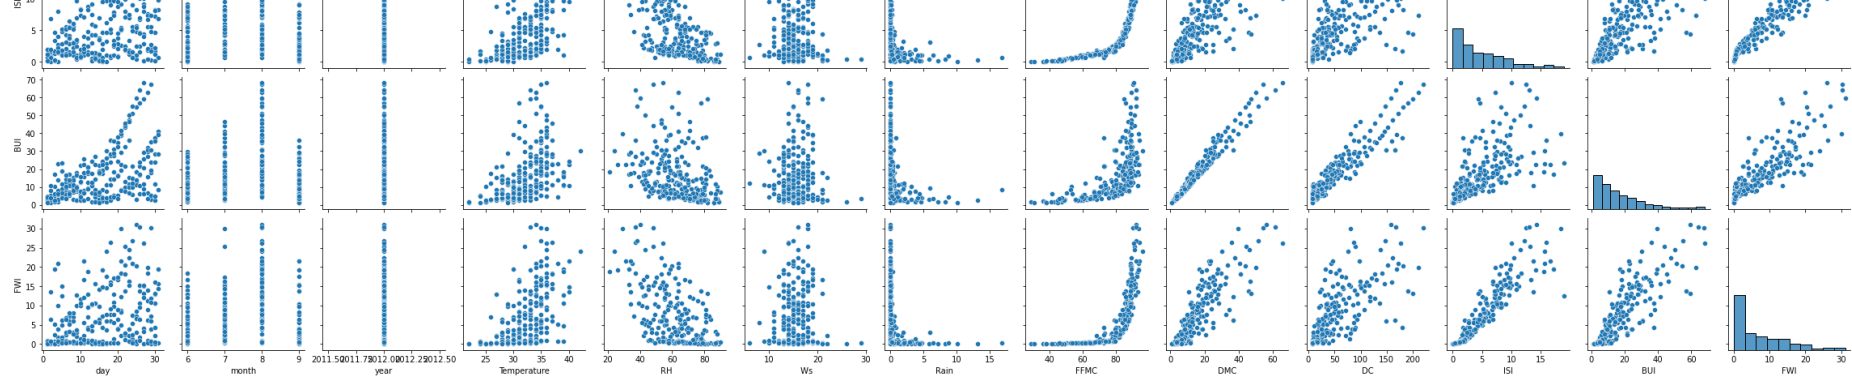
```
In [26]: 1 df['year'].unique()
```

Out[26]: array([2012], dtype=int64)

```
In [27]: 1 sns.pairplot(df)
```

```
Out[27]: <seaborn.axisgrid.PairGrid at 0x17b2055cd60>
```





```
In [28]: 1 sns.set(rc={'figure.figsize':(12,10)})
        2 sns.heatmap(df.corr(), annot=True)
```

Out[28]: <AxesSubplot:>



Observations:


1. DMC, BUI highly correlated
2. ISI, FWI highly correlated

Checking correlation with Temperature and above columns


```
In [29]: 1 df.columns
```


```
Out[29]: Index(['day', 'month', 'year', 'Temperature', ' RH', ' Ws', 'Rain ', 'FFMC',  
              'DMC', 'DC', 'ISI', 'BUI', 'FWI', 'Classes  '],  
              dtype='object')
```

```
In [30]: 1 sns.scatterplot(x='Temperature', y='DMC', data=df)
```




```
In [31]: 1 sns.scatterplot(x='Temperature', y='BUI', data=df)
```

```
Out[31]: <AxesSubplot:xlabel='Temperature', ylabel='BUI'>
```





```
In [37]: 1 sns.scatterplot(x='Temperature', y='ISI', data=df)
```

```
Out[37]: <AxesSubplot:xlabel='Temperature', ylabel='ISI'>
```



```
In [35]: 1 sns.scatterplot(x='Temperature', y='FWI', data=df)
```

```
Out[35]: <AxesSubplot:xlabel='Temperature', ylabel='FWI'>
```




Observations:

1. ISI and FWI have almost similar correlation with Temperature column
2. BUI and DMC have almost similar correlation with Temperature column

```
In [36]: 1 sns.regplot(x='Temperature', y=' RH', data=df)
```


```
Out[36]: <AxesSubplot:xlabel='Temperature', ylabel=' RH'>
```



Observation:

Temperature is inversely proportional to Relative Humidity

```
In [40]: 1 sns.regplot(x='Temperature', y='FFMC', data=df)
```




Observation:

Temperature is directly proportional to Fine Fuel Moisture Code (FFMC) index

```
In [41]: 1 sns.regplot(x='Temperature', y='FWI', data=df)
```

```
Out[41]: <AxesSubplot:xlabel='Temperature', ylabel='FWI'>
```



Observation:

Temperature is directly proportional to Fire Weather Index (FWI)

Checking Outliers

```
In [45]: 1 sns.boxplot(df['DC'])
```

```
C:\Users\pc\anaconda3\lib\site-packages\seaborn\_decorators.py:36: FutureWarning: Pass the following variable as a key word arg: x. From version 0.12, the only valid positional argument will be `data`, and passing other arguments without an explicit keyword will result in an error or misinterpretation.
  warnings.warn(
```

```
Out[45]: <AxesSubplot:xlabel='DC'>
```

Handling Categorical Variables

```
In [71]: 1 df['Classes '] = df['Classes '].str.strip()
```

```
In [72]: 1 df['Classes'].unique()
```

```
Out[72]: array(['not fire', 'fire'], dtype=object)
```

```
In [ ]: 1 replace = []
```

```
In [75]: 1 df['Classes'] = df['Classes'].replace(['not fire', 'fire'], [0,1])
```

```
In [77]: 1 df.sample(5)
```

```
Out[77]:
```

	day	month	year	Temperature	RH	Ws	Rain	FFMC	DMC	DC	ISI	BUI	FWI	Classes
83	23	8	2012	36	53	16	0.0	89.5	37.6	161.5	10.4	47.5	22.3	1
2	3	6	2012	26	82	22	13.1	47.1	2.5	7.1	0.3	2.7	0.1	0
226	13	9	2012	29	49	19	0.0	88.6	11.5	33.4	9.1	12.4	10.3	1
194	12	8	2012	39	21	17	0.4	93.0	18.4	41.5	15.5	18.4	18.8	1
66	6	8	2012	32	75	14	0.0	86.4	13.0	39.1	5.2	14.2	6.8	1

Independent and Dependent Features

```
In [80]: 1 df.head()
```

```
Out[80]:
```

	day	month	year	Temperature	RH	Ws	Rain	FFMC	DMC	DC	ISI	BUI	FWI	Classes
0	1	6	2012	29	57	18	0.0	65.7	3.4	7.6	1.3	3.4	0.5	0
1	2	6	2012	29	61	13	1.3	64.4	4.1	7.6	1.0	3.9	0.4	0
2	3	6	2012	26	82	22	13.1	47.1	2.5	7.1	0.3	2.7	0.1	0
3	4	6	2012	25	89	13	2.5	28.6	1.3	6.9	0.0	1.7	0.0	0
4	5	6	2012	27	77	16	0.0	64.8	3.0	14.2	1.2	3.9	0.5	0

```
In [81]: 1 X = df.drop(columns='Temperature')
```

```
In [82]: 1 X
```

Out[82]:

	day	month	year	RH	Ws	Rain	FFMC	DMC	DC	ISI	BUI	FWI	Classes
0	1	6	2012	57	18	0.0	65.7	3.4	7.6	1.3	3.4	0.5	0
1	2	6	2012	61	13	1.3	64.4	4.1	7.6	1.0	3.9	0.4	0
2	3	6	2012	82	22	13.1	47.1	2.5	7.1	0.3	2.7	0.1	0
3	4	6	2012	89	13	2.5	28.6	1.3	6.9	0.0	1.7	0.0	0
4	5	6	2012	77	16	0.0	64.8	3.0	14.2	1.2	3.9	0.5	0
...
239	26	9	2012	65	14	0.0	85.4	16.0	44.5	4.5	16.9	6.5	1
240	27	9	2012	87	15	4.4	41.1	6.5	8.0	0.1	6.2	0.0	0
241	28	9	2012	87	29	0.5	45.9	3.5	7.9	0.4	3.4	0.2	0
242	29	9	2012	54	18	0.1	79.7	4.3	15.2	1.7	5.1	0.7	0
243	30	9	2012	64	15	0.2	67.3	3.8	16.5	1.2	4.8	0.5	0

243 rows × 13 columns

```
In [83]: 1 y = df['Temperature']
```

```
In [84]: 1 y
```

Out[84]:

0	29
1	29
2	26
3	25
4	27
	..
239	30
240	28
241	27
242	24
243	24

Name: Temperature, Length: 243, dtype: int64

split data into training and test set

```
In [85]: 1 from sklearn.model_selection import train_test_split
```

```
In [86]: 1 X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.3, random_state=42)
```

```
In [87]: 1 X_train
```

Out[87]:

	day	month	year	RH	Ws	Rain	FFMC	DMC	DC	ISI	BUI	FWI	Classes	
	29	30	6	2012	50	14	0.0	88.7	22.9	92.8	7.2	28.3	12.9	1
	120	29	9	2012	80	16	1.8	47.4	2.9	7.7	0.3	3.0	0.1	0
	114	23	9	2012	54	11	0.5	73.7	7.9	30.4	1.2	9.6	0.7	0
	242	29	9	2012	54	18	0.1	79.7	4.3	15.2	1.7	5.1	0.7	0
	5	6	6	2012	67	14	0.0	82.6	5.8	22.2	3.1	7.0	2.5	1

	106	15	9	2012	82	15	0.4	44.9	0.9	7.3	0.2	1.4	0.0	0
	14	15	6	2012	80	17	3.1	49.4	3.0	7.4	0.4	3.0	0.1	0
	92	1	9	2012	76	17	7.2	46.0	1.3	7.5	0.2	1.8	0.1	0
	180	29	7	2012	59	16	0.0	88.1	19.5	47.2	7.4	19.5	10.9	1
	102	11	9	2012	77	21	1.8	58.5	1.9	8.4	1.1	2.4	0.3	0

170 rows × 13 columns

```
In [88]: 1 X_train.shape
```

Out[88]: (170, 13)

```
In [89]: 1 X_test
```

Out[89]:

	day	month	year	RH	Ws	Rain	FFMC	DMC	DC	ISI	BUI	FWI	Classes
24	25	6	2012	64	15	0.0	86.7	14.2	63.8	5.7	18.3	8.4	1
6	7	6	2012	54	13	0.0	88.2	9.9	30.5	6.4	10.9	7.2	1
152	1	7	2012	58	18	2.2	63.7	3.2	8.5	1.2	3.3	0.5	0
233	20	9	2012	58	13	0.2	79.5	18.7	88.0	2.1	24.4	3.8	0
239	26	9	2012	65	14	0.0	85.4	16.0	44.5	4.5	16.9	6.5	1
...
195	13	8	2012	34	16	0.2	88.3	16.9	45.1	7.5	17.5	10.5	1
104	13	9	2012	86	21	4.6	40.9	1.3	7.5	0.1	1.8	0.0	0
109	18	9	2012	49	11	0.0	89.4	9.8	33.1	6.8	11.3	7.7	1
191	9	8	2012	43	12	0.0	91.7	16.5	30.9	9.6	16.4	12.7	1
79	19	8	2012	62	19	0.0	89.4	23.2	120.9	9.7	31.3	17.2	1

73 rows × 13 columns

```
In [90]: 1 X_test.shape
```

Out[90]: (73, 13)

```
In [91]: 1 y_train
```

Out[91]: 29 33
120 26
114 32
242 24
5 31
..
106 24
14 28
92 25
180 34
102 30
Name: Temperature, Length: 170, dtype: int64

```
In [92]: 1 y_train.shape
```

```
Out[92]: (170,)
```

```
In [93]: 1 y_test
```

```
Out[93]: 24      31
          6       33
          152     28
          233     34
          239     30
          ..
          195     35
          104     25
          109     32
          191     39
          79      35
          Name: Temperature, Length: 73, dtype: int64
```

```
In [94]: 1 y_test.shape
```

```
Out[94]: (73,)
```

Feature Scaling

```
In [95]: 1 from sklearn.preprocessing import StandardScaler
```

```
In [96]: 1 scaler = StandardScaler()
```

```
In [97]: 1 X_train = scaler.fit_transform(X_train)
```

```
In [98]: 1 X_test = scaler.transform(X_test)
```

```
In [99]: 1 X_train
```

```
Out[99]: array([[ 1.56765151, -1.30687831,  0.          , ...,  0.75507842,
                  0.74777936,  0.90992142],
                [ 1.45605153,  1.39153439,  0.          , ..., -0.94987343,
                 -0.91039641, -1.098996   ],
                [ 0.78645164,  1.39153439,  0.          , ..., -0.50510338,
                 -0.83266942, -1.098996   ],
                ...,
                [-1.66874796,  1.39153439,  0.          , ..., -1.03074071,
                 -0.91039641, -1.098996   ],
                [ 1.45605153, -0.40740741,  0.          , ...,  0.16205169,
                  0.48868939,  0.90992142],
                [-0.55274815,  1.39153439,  0.          , ..., -0.99030707,
                 -0.88448741, -1.098996   ]])
```

```
In [100]: 1 X_test
```

```
9.09921419e-01],
 [ 1.67925149e+00, -4.07407407e-01,  0.00000000e+00,
   8.39147711e-02,  5.55707612e-01, -3.85823876e-01,
   6.82141084e-01,  1.31398887e+00,  1.93071213e+00,
   4.79526983e-01,  1.62440169e+00,  1.11050531e+00,
   9.09921419e-01],
 [ 3.40051709e-01, -1.30687831e+00,  0.00000000e+00,
  -5.24566744e-01,  1.73236786e-01, -3.40052726e-01,
   1.70390942e-01, -8.11427023e-01, -7.00737215e-01,
  -5.22858561e-01, -7.94877806e-01, -7.41987934e-01,
  -1.09899600e+00],
 [-1.22234804e+00, -1.30687831e+00,  0.00000000e+00,
  -1.86521458e-01, -5.91704866e-01, -2.94281575e-01,
  -2.58967840e-02, -6.95072139e-01, -6.68275234e-01,
  -6.86037603e-01, -7.14010525e-01, -8.06760425e-01,
  -1.09899600e+00],
 [-1.22234804e+00,  1.39153439e+00,  0.00000000e+00,
   8.27614401e-01,  1.73236786e-01, -3.85823876e-01,
   2.33483425e-01, -8.96753939e-01, -5.38427309e-01,
  -4.52924686e-01, -8.08355686e-01, -7.03124439e-01,
```

Model Training

```
In [101]: 1 from sklearn.linear_model import LinearRegression, Lasso, Ridge, ElasticNet
```

```
In [102]: 1 regression = LinearRegression()
          2 lasso = Lasso()
          3 ridge = Ridge()
          4 elasticnet = ElasticNet()
```

```
In [103]: 1 models = [regression, lasso, ridge, elasticnet]
```

```
In [105]: 1 regression.fit(X_train, y_train)
```

```
Out[105]: ▾ LinearRegression
          LinearRegression()
```

print coefficients and intercept

```
In [106]: 1 regression.coef_
```

```
Out[106]: array([-5.80896121e-01, -2.69620725e-01,  2.99760217e-15, -9.57968404e-01,
                -7.84242778e-01,  4.24923494e-01,  1.55454106e+00,  3.50343864e+00,
                 2.70599802e+00,  4.94774573e-01, -5.04327581e+00, -3.18139966e-01,
                -1.16724927e-01])
```

```
In [107]: 1 regression.intercept_
```

```
Out[107]: 32.04117647058823
```

prediction for test data

```
In [108]: 1 y_pred = regression.predict(X_test)
```



```
In [109]:
```

1	y_pred
---	--------


```
Out[109]: array([32.79641075, 35.02274903, 30.63350595, 33.12299224, 31.94283052,
                32.23833383, 31.30709182, 34.64596675, 31.96124513, 30.82888561,
                28.35059241, 39.73626564, 34.37012538, 34.46787268, 34.06057298,
                32.84166182, 32.99177473, 25.89504057, 32.81746314, 34.91043925,
                31.04033313, 28.47749154, 33.24853634, 28.98216672, 36.63268857,
                33.89017314, 33.48485866, 33.53634489, 26.95658767, 33.53120014,
                29.6252971 , 32.30899838, 32.1633019 , 32.96925253, 32.05849405,
                32.9839793 , 31.01781389, 34.31315509, 26.67234481, 20.88467142,
                34.28671884, 32.75694265, 34.19553275, 25.0935279 , 36.4189878 ,
                32.71875823, 30.86736383, 30.67971234, 33.15003874, 28.26380837,
                37.19062884, 35.21819978, 33.66102279, 34.75102604, 33.28117209,
                32.31620419, 32.31942747, 32.54349773, 31.79147197, 36.22333466,
                33.20847127, 30.11380025, 29.12870902, 36.78820367, 31.57165732,
                31.46204735, 33.1511071 , 34.41965372, 35.04570101, 24.471576 ,
                34.64175337, 35.89062008, 33.18548941])
```

Checking Assumptions

1. linear relationship between test and predicted values

```
In [111]: 1 plt.scatter(y_test, y_pred)
          2 plt.xlabel("Test Truth data")
          3 plt.ylabel("Test predicted data")
```

```
Out[111]: Text(0, 0.5, 'Test predicted data')
```



Observation

- test values are directly proportional to predicted values
- model is good fit


2. residuals should be normally distributed

```
In [112]: 1 residuals = y_test - y_pred
```

```
In [113]: 1 residuals
```

```
Out[113]: 24    -1.796411
          6    -2.022749
          152  -2.633506
          233   0.877008
          239  -1.942831
          ...
          195  -0.045701
          104   0.528424
          109  -2.641753
          191   3.109380
          79    1.814511
          Name: Temperature, Length: 73, dtype: float64
```

```
In [114]: 1 sns.displot(residuals,kind='kde');
```




Observation

- residuals are normally distributed
- model is good fit

3. Homoscedasticity

```
In [115]: 1 sns.scatterplot(x = y_pred, y=residuals)
```

```
Out[115]: <AxesSubplot:ylabel='Temperature'>
```



Observation:

- residuals are uniformly spread
- model is good fit

Performanc Matrix

```
In [116]: 1 from sklearn.metrics import mean_absolute_error, mean_squared_error
```

```
In [120]: 1 print(f"Mean Squarred Error: {mean_squared_error(y_test, y_pred)}")
```

Mean Squarred Error: 6.399727193150853

```
In [121]: 1 print(f"Root Mean Squarred Error: {np.sqrt(mean_squared_error(y_test, y_pred))}")
```

Root Mean Squarred Error: 2.5297682093723237

```
In [122]: 1 print(f"Mean Absolute Error: {mean_absolute_error(y_test, y_pred)}")
```

Mean Absolute Error: 2.0093401158864626

R Squarred and Adjusted R Squarred

```
In [123]: 1 from sklearn.metrics import r2_score
```

```
In [124]: 1 score = r2_score(y_test, y_pred)
```

```
In [125]: 1 print(f"R Squarred: {score}")
```

R Squarred: 0.47488457776767024

```
In [127]: 1 adj_rsquare = 1 - (1-score)*(len(y_test) - 1)/(len(y_test) - X_test.shape[1] - 1)
2 print(f"Adjusted R Squarred: {adj_rsquare}")
```

Adjusted R Squarred: 0.35918117964868235

Checking Performance of Model for Ridge, Lasso and ElasticNet

```
In [135]: 1 def evaluation(model, X_train, X_test, y_train, y_test) :
2     model.fit(X_train, y_train)
3     print(f"Coefficient: \n{model.coef_}")
4     print(f"Intercept: {model.intercept_}")
5     y_pred = model.predict(X_test)
6     print(f"\n{' '*25} Evaluation for {model} model {' '*25}\n")
7     print(f"Mean Squarred Error: {mean_squared_error(y_test, y_pred)}")
8     print(f"Root Mean Squarred Error: {np.sqrt(mean_squared_error(y_test, y_pred))}")
9     print(f"Mean Absolute Error: {mean_absolute_error(y_test, y_pred)}")
10
11     score = r2_score(y_test, y_pred)
12     print(f"\nR Squarred: {score}")
13     adj_rsquare = 1 - (1-score)*(len(y_test) - 1)/(len(y_test) - X_test.shape[1] - 1)
14     print(f"Adjusted R Squarred: {adj_rsquare}")
15     print()
```

In [136]:

```
1 for model in models :  
2     evaluation(model, X_train, X_test,y_train, y_test)
```

Coefficient:

```
[-5.80896121e-01 -2.69620725e-01  2.99760217e-15 -9.57968404e-01  
 -7.84242778e-01  4.24923494e-01  1.55454106e+00  3.50343864e+00  
  2.70599802e+00  4.94774573e-01 -5.04327581e+00 -3.18139966e-01  
 -1.16724927e-01]
```

Intercept: 32.04117647058823

***** Evaluation for LinearRegression() model *****

Mean Squarred Error: 6.399727193150853

Root Mean Squarred Error: 2.5297682093723237

Mean Absolute Error: 2.0093401158864626

R Squarred: 0.47488457776767024

Adjusted R Squarred: 0.35918117964868235

Coefficient:

```
[ 0.          -0.          0.          -0.65210487 -0.          -0.  
 1.08755359  0.          0.          0.          0.          0.  
 0.          ]
```

Intercept: 32.04117647058823

***** Evaluation for Lasso() model *****

Mean Squarred Error: 6.6495770305009945

Root Mean Squarred Error: 2.5786773800731635

Mean Absolute Error: 2.123185084516282

R Squarred: 0.4543837034530257

Adjusted R Squarred: 0.3341631635358957

Coefficient:

```
[-0.57163213 -0.24486263  0.          -0.97026732 -0.7811602  0.38521533  
  1.48287006  0.45825186  1.03142175  0.51435613 -0.46476688 -0.31517238  
 -0.1118449 ]
```

Intercept: 32.04117647058823

***** Evaluation for Ridge() model *****

Mean Squarred Error: 5.93403273310624

Root Mean Squarred Error: 2.435987014149755

Mean Absolute Error: 1.9296082189973855

R Squarred: 0.5130961039213632

Adjusted R Squared: 0.4058121946159008

Coefficient:

```
[-0.          -0.          0.          -0.62675812 -0.22498622 -0.
 0.69959647  0.09593062  0.          0.26773045  0.06300842  0.17640915
 0.117818   ]
```

Intercept: 32.04117647058823

***** Evaluation for ElasticNet() model *****

Mean Squared Error: 6.184150209485093

Root Mean Squared Error: 2.4867951683814034

Mean Absolute Error: 2.007985784009525

R Squared: 0.4925732690797576

Adjusted R Squared: 0.3807673792159755

Feature Selection

```
In [138]: 1 df.columns
```

```
Out[138]: Index(['day', 'month', 'year', 'Temperature', ' RH', ' Ws', 'Rain ', 'FFMC',
                'DMC', 'DC', 'ISI', 'BUI', 'FWI', 'Classes  '],
                dtype='object')
```

1. Variance Threshold

```
In [140]: 1 from sklearn.feature_selection import VarianceThreshold
```

```
In [157]: 1 variance_threshold = VarianceThreshold(threshold=0)
          2 variance_threshold.fit(df)
```

```
Out[157]: ▾      VarianceThreshold
          VarianceThreshold(threshold=0)
```



```
In [158]: 1 variance_threshold.get_support()
```

```
Out[158]: array([ True,  True, False,  True,  True,  True,  True,  True,  True,
        True,  True,  True,  True,  True])
```

```
In [159]: 1 df.columns[variance_threshold.get_support()]
```

```
Out[159]: Index(['day', 'month', 'Temperature', ' RH', ' Ws', 'Rain ', 'FFMC', 'DMC',
        'DC', 'ISI', 'BUI', 'FWI', 'Classes  '],
        dtype='object')
```

finding columns having zero standard deviation

```
In [147]: 1 constant_columns = [column for column in df.columns if column not in df.columns[variance_threshold.get_support()]]
```

```
In [148]: 1 print(len(constant_columns))
```

```
1
```

```
In [149]: 1 for feature in constant_columns :
          2     print(feature)
```

```
year
```

```
In [167]: 1 df = df.drop(constant_columns,axis=1)
```

In [168]:

1	df
---	----

Out[168]:

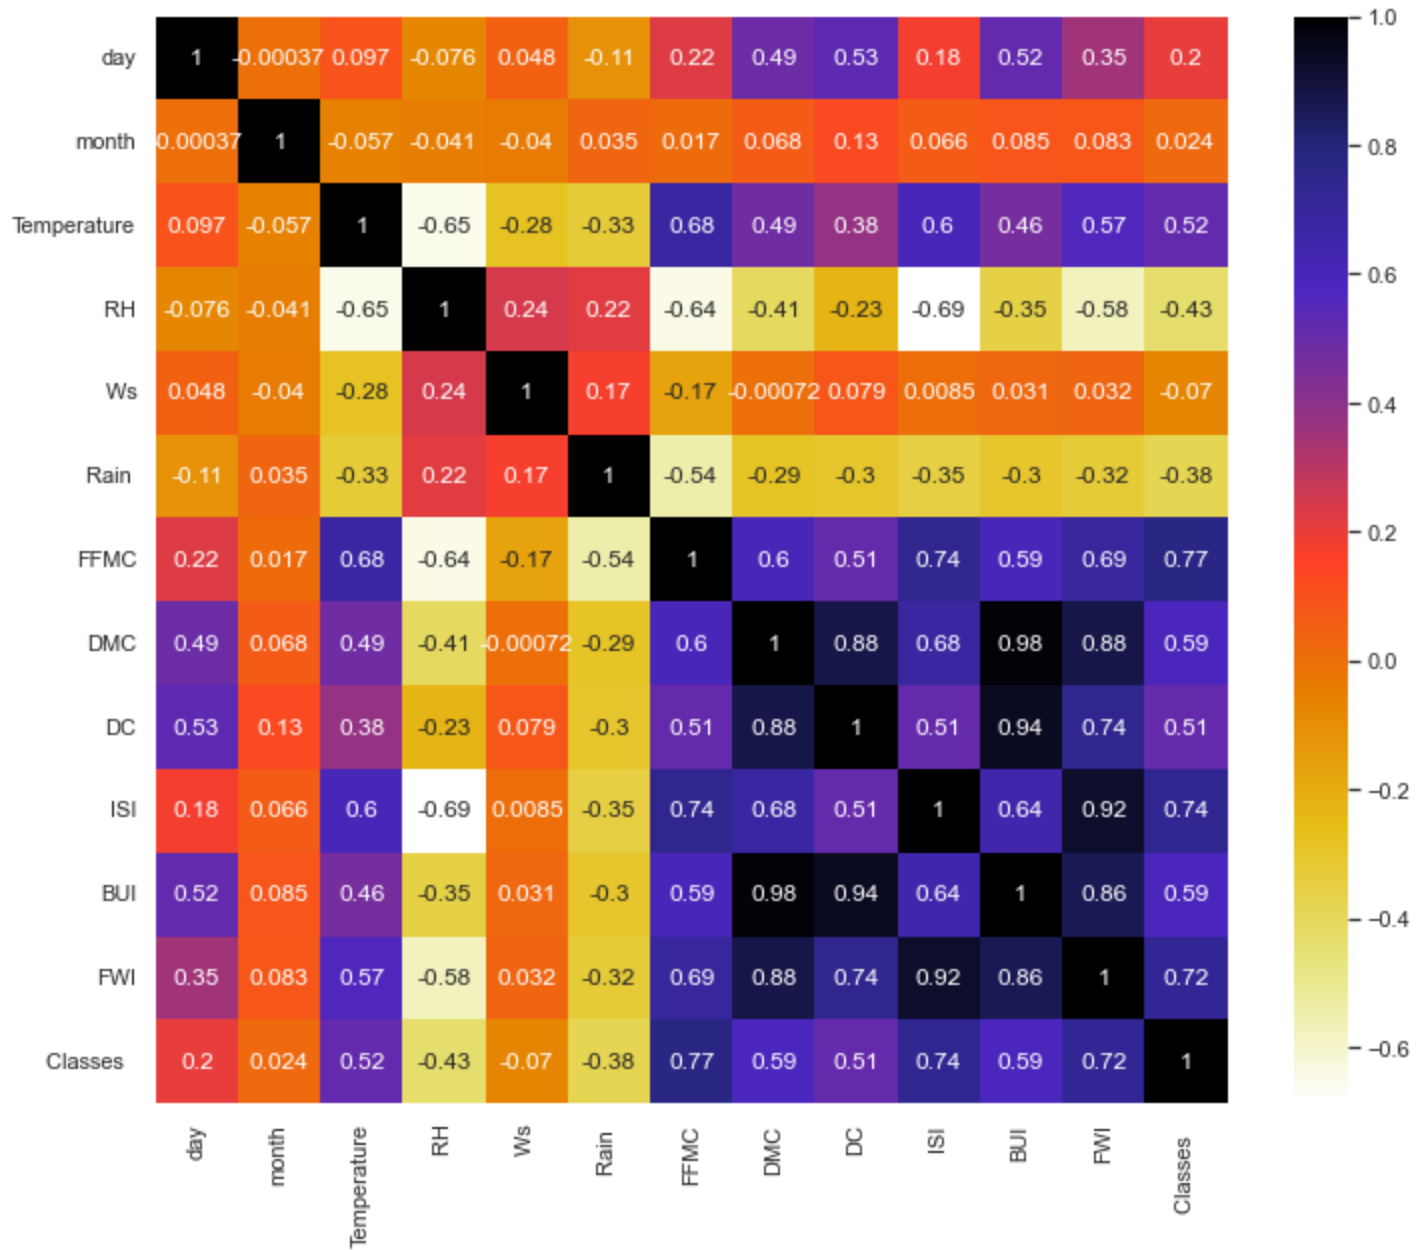
	day	month	Temperature	RH	Ws	Rain	FFMC	DMC	DC	ISI	BUI	FWI	Classes
0	1	6	29	57	18	0.0	65.7	3.4	7.6	1.3	3.4	0.5	0
1	2	6	29	61	13	1.3	64.4	4.1	7.6	1.0	3.9	0.4	0
2	3	6	26	82	22	13.1	47.1	2.5	7.1	0.3	2.7	0.1	0
3	4	6	25	89	13	2.5	28.6	1.3	6.9	0.0	1.7	0.0	0
4	5	6	27	77	16	0.0	64.8	3.0	14.2	1.2	3.9	0.5	0
...
239	26	9	30	65	14	0.0	85.4	16.0	44.5	4.5	16.9	6.5	1
240	27	9	28	87	15	4.4	41.1	6.5	8.0	0.1	6.2	0.0	0
241	28	9	27	87	29	0.5	45.9	3.5	7.9	0.4	3.4	0.2	0
242	29	9	24	54	18	0.1	79.7	4.3	15.2	1.7	5.1	0.7	0
243	30	9	24	64	15	0.2	67.3	3.8	16.5	1.2	4.8	0.5	0

243 rows × 13 columns

2. Correlation

```
In [169]: 1 sns.set(rc={'figure.figsize':(12,10)})
          2 sns.heatmap(df.corr(), annot=True, cmap=plt.cm.CMRmap_r)
```

Out[169]: <AxesSubplot:>



```
In [170]: 1 # with the following function we can select highly correlated features
2 # it will remove the first feature that is correlated with anything other feature
3
4 def correlation(dataset, threshold):
5     col_corr = set() # Set of all the names of correlated columns
6     corr_matrix = df.corr()
7     for i in range(len(corr_matrix.columns)):
8         for j in range(i):
9             if abs(corr_matrix.iloc[i, j]) > threshold: # we are interested in absolute coeff value
10                 colname = corr_matrix.columns[i] # getting the name of column
11                 col_corr.add(colname)
12     return col_corr
```

```
In [171]: 1 corr_features = correlation(df, 0.85)
2 len(set(corr_features))
```

Out[171]: 3

```
In [172]: 1 corr_features
```

Out[172]: {'BUI', 'DC', 'FWI'}

```
In [176]: 1 df = df.drop(corr_features,axis=1)
```

In [177]:

```
1 df
```

Out[177]:

	day	month	Temperature	RH	Ws	Rain	FFMC	DMC	ISI	Classes
0	1	6	29	57	18	0.0	65.7	3.4	1.3	0
1	2	6	29	61	13	1.3	64.4	4.1	1.0	0
2	3	6	26	82	22	13.1	47.1	2.5	0.3	0
3	4	6	25	89	13	2.5	28.6	1.3	0.0	0
4	5	6	27	77	16	0.0	64.8	3.0	1.2	0
...
239	26	9	30	65	14	0.0	85.4	16.0	4.5	1
240	27	9	28	87	15	4.4	41.1	6.5	0.1	0
241	28	9	27	87	29	0.5	45.9	3.5	0.4	0
242	29	9	24	54	18	0.1	79.7	4.3	1.7	0
243	30	9	24	64	15	0.2	67.3	3.8	1.2	0

243 rows × 10 columns

In [179]:

```
1 df = df.drop(columns=['day', 'month'], axis=1)
```

```
In [180]: 1 df
```

Out[180]:

	Temperature	RH	Ws	Rain	FFMC	DMC	ISI	Classes
0	29	57	18	0.0	65.7	3.4	1.3	0
1	29	61	13	1.3	64.4	4.1	1.0	0
2	26	82	22	13.1	47.1	2.5	0.3	0
3	25	89	13	2.5	28.6	1.3	0.0	0
4	27	77	16	0.0	64.8	3.0	1.2	0
...
239	30	65	14	0.0	85.4	16.0	4.5	1
240	28	87	15	4.4	41.1	6.5	0.1	0
241	27	87	29	0.5	45.9	3.5	0.4	0
242	24	54	18	0.1	79.7	4.3	1.7	0
243	24	64	15	0.2	67.3	3.8	1.2	0

243 rows × 8 columns

repeating same procedure to find model performance

```
In [181]: 1 X = df.drop(columns='Temperature')
```

```
In [182]: 1 X
```

Out[182]:

	RH	Ws	Rain	FFMC	DMC	ISI	Classes
0	57	18	0.0	65.7	3.4	1.3	0
1	61	13	1.3	64.4	4.1	1.0	0
2	82	22	13.1	47.1	2.5	0.3	0
3	89	13	2.5	28.6	1.3	0.0	0
4	77	16	0.0	64.8	3.0	1.2	0
...
239	65	14	0.0	85.4	16.0	4.5	1
240	87	15	4.4	41.1	6.5	0.1	0
241	87	29	0.5	45.9	3.5	0.4	0
242	54	18	0.1	79.7	4.3	1.7	0
243	64	15	0.2	67.3	3.8	1.2	0

243 rows × 7 columns

```
In [183]: 1 y = df['Temperature']
```

```
In [184]: 1 y
```

Out[184]:

0	29
1	29
2	26
3	25
4	27
	..
239	30
240	28
241	27
242	24
243	24

Name: Temperature, Length: 243, dtype: int64

split data into training and test set

split data into training and test set

```
In [185]: 1 X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.3, random_state=42)
```

```
In [186]: 1 X_train
```

Out[186]:

	RH	Ws	Rain	FFMC	DMC	ISI	Classes
29	50	14	0.0	88.7	22.9	7.2	1
120	80	16	1.8	47.4	2.9	0.3	0
114	54	11	0.5	73.7	7.9	1.2	0
242	54	18	0.1	79.7	4.3	1.7	0
5	67	14	0.0	82.6	5.8	3.1	1
...
106	82	15	0.4	44.9	0.9	0.2	0
14	80	17	3.1	49.4	3.0	0.4	0
92	76	17	7.2	46.0	1.3	0.2	0
180	59	16	0.0	88.1	19.5	7.4	1
102	77	21	1.8	58.5	1.9	1.1	0

170 rows × 7 columns

```
In [187]: 1 X_train.shape
```

Out[187]: (170, 7)


```
In [188]: 1 X_test
```

Out[188]:

	RH	Ws	Rain	FFMC	DMC	ISI	Classes
24	64	15	0.0	86.7	14.2	5.7	1
6	54	13	0.0	88.2	9.9	6.4	1
152	58	18	2.2	63.7	3.2	1.2	0
233	58	13	0.2	79.5	18.7	2.1	0
239	65	14	0.0	85.4	16.0	4.5	1
...
195	34	16	0.2	88.3	16.9	7.5	1
104	86	21	4.6	40.9	1.3	0.1	0
109	49	11	0.0	89.4	9.8	6.8	1
191	43	12	0.0	91.7	16.5	9.6	1
79	62	19	0.0	89.4	23.2	9.7	1

73 rows × 7 columns

```
In [189]: 1 X_test.shape
```

Out[189]: (73, 7)

```
In [190]: 1 y_train
```

Out[190]: 29 33
120 26
114 32
242 24
5 31
..
106 24
14 28
92 25
180 34
102 30
Name: Temperature, Length: 170, dtype: int64

```
In [191]: 1 y_train.shape
```

```
Out[191]: (170,)
```

```
In [192]: 1 y_test
```

```
Out[192]: 24      31
          6       33
          152     28
          233     34
          239     30
          ..
          195     35
          104     25
          109     32
          191     39
          79      35
          Name: Temperature, Length: 73, dtype: int64
```

```
In [193]: 1 y_test.shape
```

```
Out[193]: (73,)
```

Feature Scaling

```
In [194]: 1 from sklearn.preprocessing import StandardScaler
```

```
In [195]: 1 scaler = StandardScaler()
```

```
In [196]: 1 X_train = scaler.fit_transform(X_train)
```

```
In [197]: 1 X_test = scaler.transform(X_test)
```

```
In [198]: 1 X_train
```

```
Out[198]: array([[ -0.86261203, -0.59170487, -0.38582388, ...,  0.61585956,
        0.57277215,  0.90992142],
       [ 1.16565969,  0.17323679,  0.43805684, ..., -0.9355389 ,
       -1.03570698, -1.098996  ],
       [-0.5921758 , -1.73911734, -0.15696812, ..., -0.54768929,
       -0.82590535, -1.098996  ],
       ...,
       [ 0.89522346,  0.55570761,  2.90969898, ..., -1.05965078,
       -1.05901827, -1.098996  ],
       [-0.25413052,  0.17323679, -0.38582388, ...,  0.35212182,
        0.61939473,  0.90992142],
       [ 0.96283252,  2.08559091,  0.43805684, ..., -1.01310882,
       -0.84921665, -1.098996  ]])
```

```
In [199]: 1 X_test
```

```
[ 1.65057158e-02, -2.09234040e-01, -3.85823876e-01,
  6.68120532e-01,  3.13336859e-01,  2.69725358e-01,
  9.09921419e-01],
[ 1.57131403e+00, -2.09234040e-01,  4.23706235e+00,
 -3.29268536e+00, -1.10619273e+00, -1.10564085e+00,
 -1.09899600e+00],
[-3.21739572e-01,  9.38178437e-01, -3.85823876e-01,
 7.38223292e-01, -1.05540725e-01,  7.59262484e-01,
 9.09921419e-01],
[ 1.43609592e+00,  9.38178437e-01, -3.85823876e-01,
 4.43791703e-01, -1.13297717e-01, -5.66327264e-02,
 9.09921419e-01],
[-3.21739572e-01, -1.35664652e+00,  1.49079330e+00,
 -7.97027136e-01, -8.50211985e-01, -8.72527937e-01,
 -1.09899600e+00],
[ 8.39147711e-02,  9.38178437e-01, -3.85823876e-01,
 6.54099981e-01,  2.20252952e-01,  4.56215692e-01,
 9.09921419e-01],
[ 8.39147711e-02,  5.55707612e-01, -3.85823876e-01,
 6.82141084e-01,  1.31398887e+00,  4.79526983e-01,
```

Model Training

```
In [200]: 1 from sklearn.linear_model import LinearRegression, Lasso, Ridge, ElasticNet
```

```
In [201]: 1 regression = LinearRegression()
          2 lasso = Lasso()
          3 ridge = Ridge()
          4 elasticnet = ElasticNet()
```

```
In [202]: 1 models = [regression, lasso, ridge, elasticnet]
```

```
In [203]: 1 for model in models :
          2     evaluation(model, X_train, X_test,y_train, y_test)
```

Intercept: 32.04117647058823

***** Evaluation for Ridge() model *****

Mean Squarred Error: 5.7797092049513905

Root Mean Squarred Error: 2.4041025778762832

Mean Absolute Error: 1.9559622622376804

R Squarred: 0.5257587787825893

Adjusted R Squarred: 0.47468664726686804

Coefficient:

[-0.643625 -0.21571403 -0. 0.71828381 0.18822364 0.32893064
 0.14715384]

Intercept: 32.04117647058823

***** Evaluation for ElasticNet() model *****

Mean Squarred Error: 6.186738102315899

Root Mean Squarred Error: 2.4872154408550202

```
In [ ]: 1
```